

Business Case 1: To prepare a complete Data Analysis Report using Flight Price Prediction Dataset

Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')
```

Importing Dataset

```
In [2]: df = pd.read_excel('Flight_Fare.xlsx')
df
```

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_Stops	Additional_Info	Price
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	22:20	01:10 22 Mar	2h 50m	non-stop	No info	3897
1	Air India	2019-05-01	Kolkata	Banglore	CCU → IXR → BBI → BLR	05:50	13:15	7h 25m	2 stops	No info	7662
2	Jet Airways	2019-06-09	Delhi	Cochin	DEL → LKO → BOM → COK	09:25	04:25 10 Jun	19h	2 stops	No info	13882
3	IndiGo	2019-05-12	Kolkata	Banglore	CCU → NAG → BLR	18:05	23:30	5h 25m	1 stop	No info	6218
4	IndiGo	2019-03-01	Banglore	New Delhi	BLR → NAG → DEL	16:50	21:35	4h 45m	1 stop	No info	13302
...
10678	Air Asia	2019-04-09	Kolkata	Banglore	CCU → BLR	19:55	22:25	2h 30m	non-stop	No info	4107
10679	Air India	2019-04-27	Kolkata	Banglore	CCU → BLR	20:45	23:20	2h 35m	non-stop	No info	4145
10680	Jet Airways	2019-04-27	Banglore	Delhi	BLR → DEL	08:20	11:20	3h	non-stop	No info	7229
10681	Vistara	2019-03-01	Banglore	New Delhi	BLR → DEL	11:30	14:10	2h 40m	non-stop	No info	12648
10682	Air India	2019-05-09	Delhi	Cochin	DEL → GOI → BOM → COK	10:55	19:15	8h 20m	2 stops	No info	11753

10683 rows × 11 columns

```
In [3]: # Copying data set to keep as Back-up
```

```
df_Backup = df.copy()
```

Domain Analysis

1. **Airline:** This column will have all the types of airlines like Indigo, Jet Airways, Air India, and many more.
2. **Date_of_Journey:** This column will let us know about the date on which the passenger's journey will start.
3. **Source:** This column holds the name of the place from where the passenger's journey will start.
4. **Destination:** This column holds the name of the place to where passengers wanted to travel.
5. **Route:** This column tells what is the route through which passengers have opted to travel from his/her source to their destination.
6. **Dep_Time:** Departure time is when the flight will depart from the source.
7. **Arrival_Time:** Arrival time is when the passenger will reach his/her destination.
8. **Duration:** Duration is the whole period that a flight will take to complete its journey from source to destination.
9. **Total_Stops:** This will let us know in how many places flights will stop there for the flight in the whole journey.
10. **Additional_Info:** In this column, we will get information about food, kind of food, and other amenities.
11. **Price:** Price of the flight for a complete journey including all the expenses before onboarding.

Basic Checks and Statistical Analysis

In [4]: df

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_Stops	Additional_Info	Price
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	22:20	01:10 22 Mar	2h 50m	non-stop	No info	3897
1	Air India	2019-05-01	Kolkata	Banglore	CCU → IXR → BBI → BLR	05:50	13:15	7h 25m	2 stops	No info	7662
2	Jet Airways	2019-06-09	Delhi	Cochin	DEL → LKO → BOM → COK	09:25	04:25 10 Jun	19h	2 stops	No info	13882
3	IndiGo	2019-05-12	Kolkata	Banglore	CCU → NAG → BLR	18:05	23:30	5h 25m	1 stop	No info	6218
4	IndiGo	2019-03-01	Banglore	New Delhi	BLR → NAG → DEL	16:50	21:35	4h 45m	1 stop	No info	13302
...
10678	Air Asia	2019-04-09	Kolkata	Banglore	CCU → BLR	19:55	22:25	2h 30m	non-stop	No info	4107
10679	Air India	2019-04-27	Kolkata	Banglore	CCU → BLR	20:45	23:20	2h 35m	non-stop	No info	4145
10680	Jet Airways	2019-04-27	Banglore	Delhi	BLR → DEL	08:20	11:20	3h	non-stop	No info	7229
10681	Vistara	2019-03-01	Banglore	New Delhi	BLR → DEL	11:30	14:10	2h 40m	non-stop	No info	12648
10682	Air India	2019-05-09	Delhi	Cochin	DEL → GOI → BOM → COK	10:55	19:15	8h 20m	2 stops	No info	11753

10683 rows × 11 columns

In [5]: df.shape

Out[5]: (10683, 11)

Observation

There are 11 features (columns) and 10683 records (rows) in the Flight Fare data set.

In [6]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Airline          10683 non-null   object  
 1   Date_of_Journey  10683 non-null   datetime64[ns]
 2   Source           10683 non-null   object  
 3   Destination      10683 non-null   object  
 4   Route            10682 non-null   object  
 5   Dep_Time         10683 non-null   object  
 6   Arrival_Time     10683 non-null   object  
 7   Duration         10683 non-null   object  
 8   Total_Stops      10682 non-null   object  
 9   Additional_Info  10683 non-null   object  
 10  Price            10683 non-null   int64  
dtypes: datetime64[ns](1), int64(1), object(9)
memory usage: 918.2+ KB
```

Observation:

- There are 10 features belonging to Object data type.
- Only the Target variable 'Price' is Integer type.
- There is only 1 missing record present in 'Route' column and 1 in 'Total_Stops' column.

In [7]: #statistical analysis for Numerical Data

df.describe()

	Date_of_Journey	Price
count	10683	10683.000000
mean	2019-05-04 19:56:41.853412096	9087.064121
min	2019-03-01 00:00:00	1759.000000
25%	2019-03-27 00:00:00	5277.000000
50%	2019-05-15 00:00:00	8372.000000
75%	2019-06-06 00:00:00	12373.000000
max	2019-06-27 00:00:00	79512.000000
std	Nan	4611.359167

Observation

- No constant Data
- No corrupt Data

In [8]: # Statistical Analysis for Categorical Data

df.describe(include = 'O')

	Airline	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_Stops	Additional_Info
count	10683	10683	10683	10682	10683	10683	10683	10682	10683
unique	12	5	6	128	222	1343	368	5	10
top	Jet Airways	Delhi	Cochin	DEL → BOM → COK	18:55	19:00	2h 50m	1 stop	No info
freq	3849	4537	4537	2376	233	423	550	5625	8345

Observation

- No constant Data

Exploratory Data Analysis

Univariate Analysis - Sweetviz Report

In [9]: # Installation

```
!pip install sweetviz
```

```
Requirement already satisfied: sweetviz in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (2.3.1)
Requirement already satisfied: pandas!=1.0.0,!>=1.0.1,!>=1.0.2,>=0.25.3 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from sweetviz) (2.3.3)
Requirement already satisfied: numpy>=1.16.0 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from sweetviz) (2.3.3)
Requirement already satisfied: matplotlib>=3.1.3 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from sweetviz) (3.10.6)
Requirement already satisfied: tqdm>=4.43.0 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from sweetviz) (4.67.1)
Requirement already satisfied: scipy>=1.3.2 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from sweetviz) (1.16.2)
Requirement already satisfied: jinja2>=2.11.1 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from sweetviz) (3.1.6)
Requirement already satisfied: importlib-resources>=1.2.0 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from sweetviz) (6.5.2)
Requirement already satisfied: MarkupSafe>=2.0 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from jinja2>=2.11.1->sweetviz) (3.0.2)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.1.3->sweetviz) (1.3.3)
Requirement already satisfied: cycler>=0.10 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.1.3->sweetviz) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.1.3->sweetviz) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.1.3->sweetviz) (1.4.9)
Requirement already satisfied: packaging>=20.0 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.1.3->sweetviz) (25.0)
Requirement already satisfied: pillow>=8 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.1.3->sweetviz) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.1.3->sweetviz) (3.2.5)
Requirement already satisfied: python-dateutil>=2.7 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from matplotlib>=3.1.3->sweetviz) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from pandas!=1.0.0,!>=1.0.1,!>=1.0.2,>=0.25.3->sweetviz) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from pandas!=1.0.0,!>=1.0.1,!>=1.0.2,>=0.25.3->sweetviz) (2025.2)
Requirement already satisfied: six>=1.5 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from python-dateutil>=2.7->matplotlib>=3.1.3->sweetviz) (1.17.0)
Requirement already satisfied: colorama in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from tqdm>=4.43.0->sweetviz) (0.4.6)
[notice] A new release of pip is available: 25.2 -> 25.3
[notice] To update, run: python.exe -m pip install --upgrade pip
```

In [10]: import sweetviz as sv

```
report = sv.analyze(df)
report.show_html()
```

Feature: Price | [92%] 00:02 -> (00:00 left)

```
AttributeError                                     Traceback (most recent call last)
Cell In[10], line 3
      1 import sweetviz as sv
----> 3 report = sv.analyze(df)
      4 report.show_html()

File ~\AppData\Local\Programs\Python\Python313\Lib\site-packages\sweetviz\sv_public.py:12, in analyze(source, target_feat, feat_cfg, pairwise_analysis)
     8 def analyze(source: Union[pd.DataFrame, Tuple[pd.DataFrame, str]], 
     9             target_feat: str = None,
    10            feat_cfg: FeatureConfig = None,
    11            pairwise_analysis: str = 'auto'):
--> 12    report = sweetviz.DataframeReport(source, target_feat, None,
    13                                         pairwise_analysis, feat_cfg)
    14    return report

File ~\AppData\Local\Programs\Python\Python313\Lib\site-packages\sweetviz\dataframe_report.py:277, in DataframeReport.__init__(self, source, target_feature_name, compare, pairwise_analysis, fc, verbosity)
274 for f in features_to_process:
275     # start = time.perf_counter()
276     self.progress_bar.set_description_str(f"Feature: {f.source.name}")
--> 277     self._features[f.source.name] = sa.analyze_feature_to_dictionary(f)
278     self.progress_bar.update(1)
279     # print(f"DONE FEATURE-----> {f.source.name}")
280     #     f" {(time.perf_counter() - start):.2f} {self._features[f.source.name]['type']}")
281 # self.progress_bar.set_description_str('[FEATURES DONE]')
282 # self.progress_bar.close()
283
284 # Wrap up summary

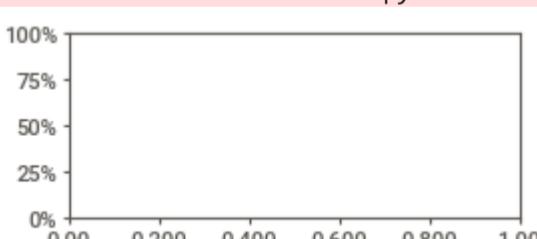
File ~\AppData\Local\Programs\Python\Python313\Lib\site-packages\sweetviz\series_analyzer.py:142, in analyze_feature_to_dictionary(to_process)
140 # Perform full analysis on source/compare/target
141 if returned_feature_dict["type"] == FeatureType.TYPE_NUM:
--> 142     sweetviz.series_analyzer_numeric.analyze(to_process, returned_feature_dict)
143 elif returned_feature_dict["type"] == FeatureType.TYPE_CAT:
144     sweetviz.series_analyzer_cat.analyze(to_process, returned_feature_dict)

File ~\AppData\Local\Programs\Python\Python313\Lib\site-packages\sweetviz\series_analyzer_numeric.py:102, in analyze(to_process, feature_dict)
 98     do_stats_numeric(to_process.compare, compare_dict)
100 do_detail_numeric(to_process.source, to_process.source_counts, to_process.compare_counts, feature_dict)
--> 102 feature_dict["minigraph"] = GraphNumeric(      , to_process)
103 feature_dict["detail_graphs"] = list()
104 for num_bins in [0, 5, 15, 30]:


File ~\AppData\Local\Programs\Python\Python313\Lib\site-packages\sweetviz\graph_numeric.py:71, in GraphNumeric.__init__(self, which_graph, to_process)
   67     normalizing_weights = norm_source
   69 gap_percent = config["Graphs"].getfloat("summary_graph_categorical_gap")
--> 71 warnings.filterwarnings('ignore', category=np.VisibleDeprecationWarning)
   72 self.hist_specs = axs.hist(plot_data, weights = normalizing_weights, bins=self.num_bins, \
   73                           rwidth = (100.0 - gap_percent) / 100.0)
   74 warnings.filterwarnings('once', category=np.VisibleDeprecationWarning)

File ~\AppData\Local\Programs\Python\Python313\Lib\site-packages\numpy\__init__.py:808, in __getattr__(attr)
 805     import numpy.char as char
 806     return char.chararray
--> 808 raise AttributeError(f"module {__name__}!r} has no attribute {attr!r}")


AttributeError: module 'numpy' has no attribute 'VisibleDeprecationWarning'
```



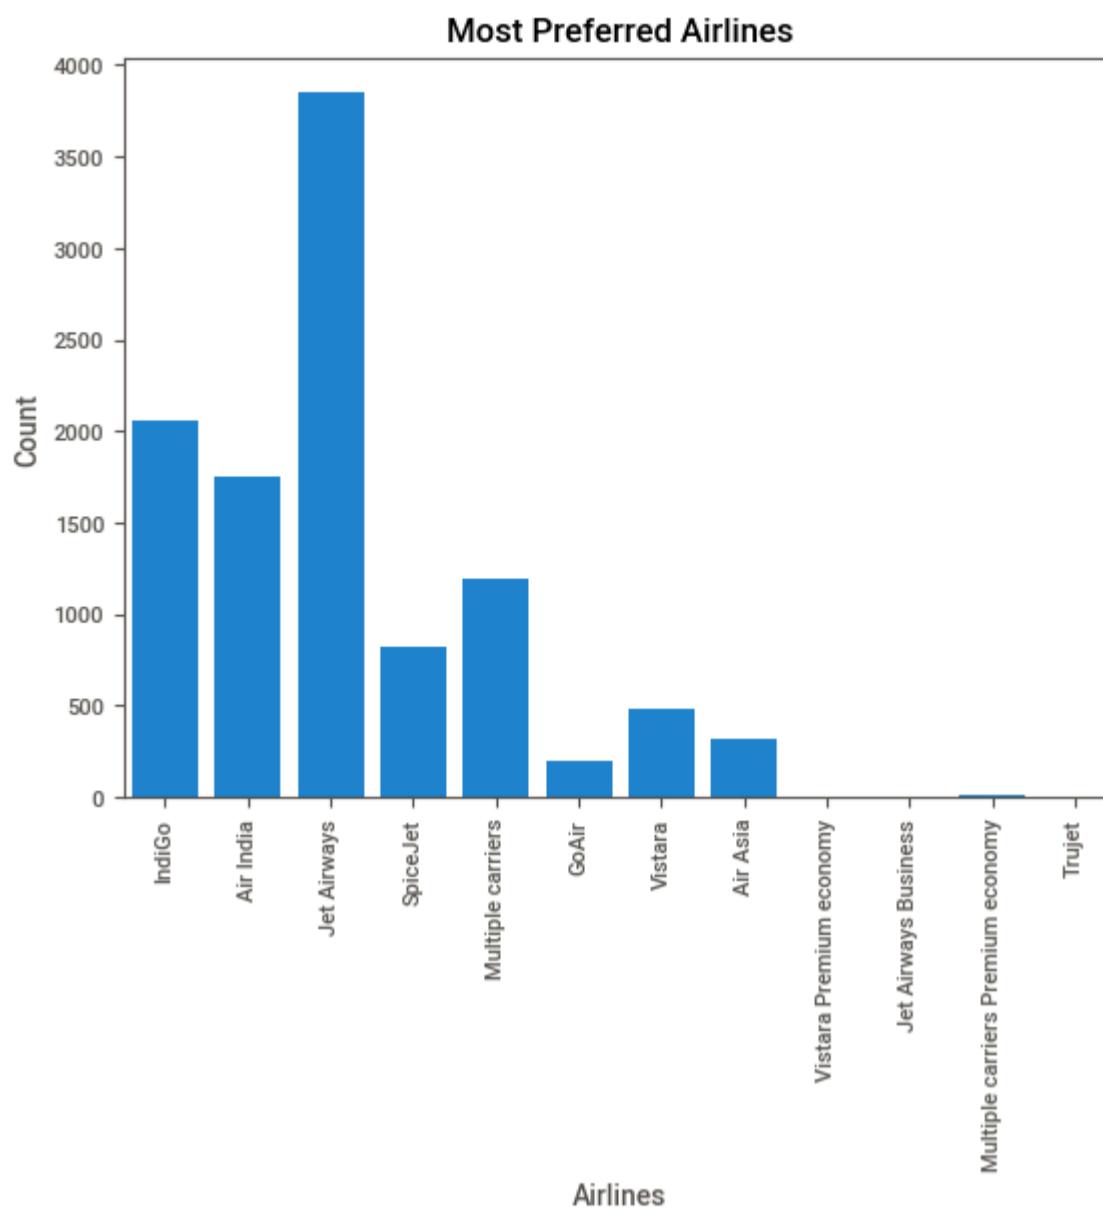
Count plot

MOST PREFERRED AIRLINE

```
In [11]: sns.countplot(x = 'Airline', data = df)

plt.title('Most Preferred Airlines', fontsize = 12, color = 'black')
plt.xlabel('Airlines', fontsize = 10)
plt.ylabel('Count', fontsize = 10)
plt.xticks(rotation = 90)

plt.show()
```



Observations:

- Jet Airways is the most preferred airline with approximately 3800 counts.
- Second most preferred Airline is Indigo with nearly 2100 counts.
- Jet Airways Business, Vistara Premium Economy, Multiple Carriers Premium economy and Trujet are the least preferred airlines.

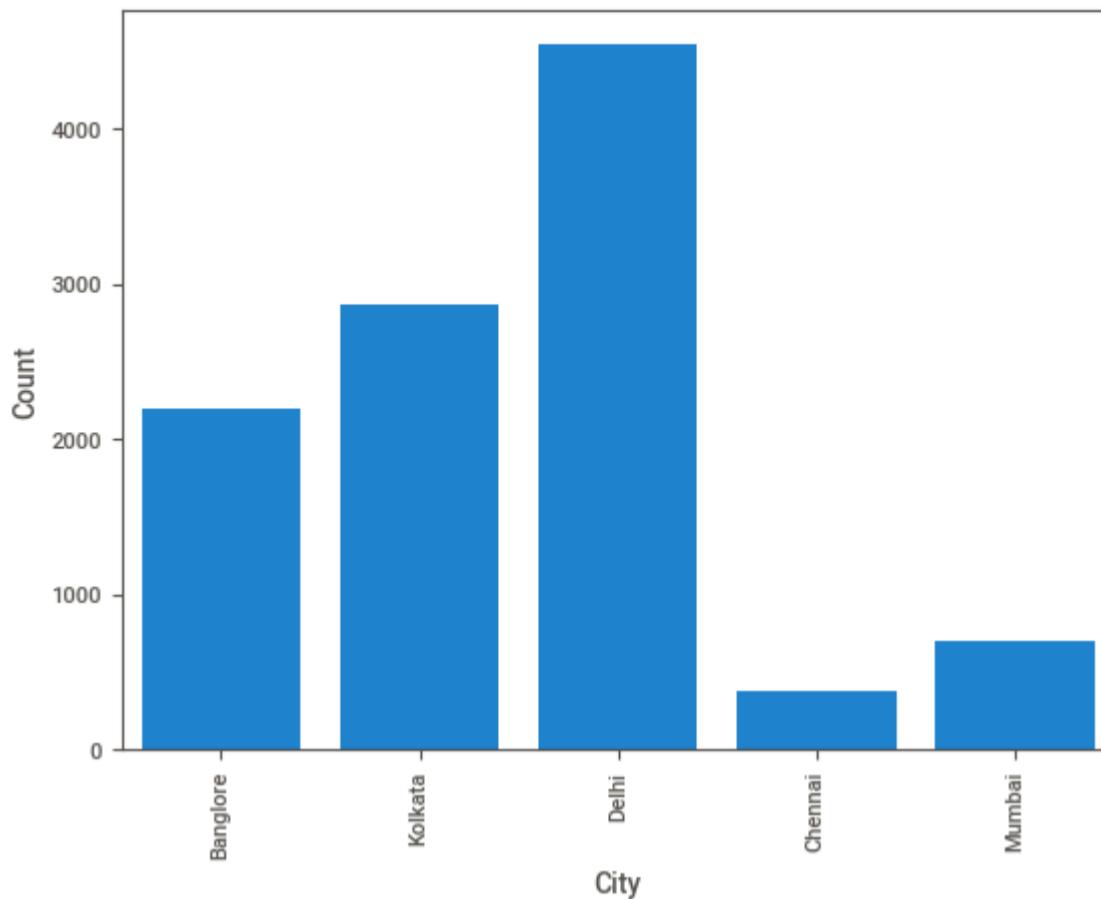
MOST PREFERRED SOURCE

```
In [12]: sns.countplot(x = 'Source', data = df)

plt.title('Most Preferred Source', fontsize = 12, color = 'black')
plt.xlabel('City', fontsize = 10)
plt.ylabel('Count', fontsize = 10)
plt.xticks(rotation = 90)
```

```
Out[12]: ([0, 1, 2, 3, 4],
[Text(0, 0, 'Banglore'),
 Text(1, 0, 'Kolkata'),
 Text(2, 0, 'Delhi'),
 Text(3, 0, 'Chennai'),
 Text(4, 0, 'Mumbai')])
```

Most Preferred Source



Observations:

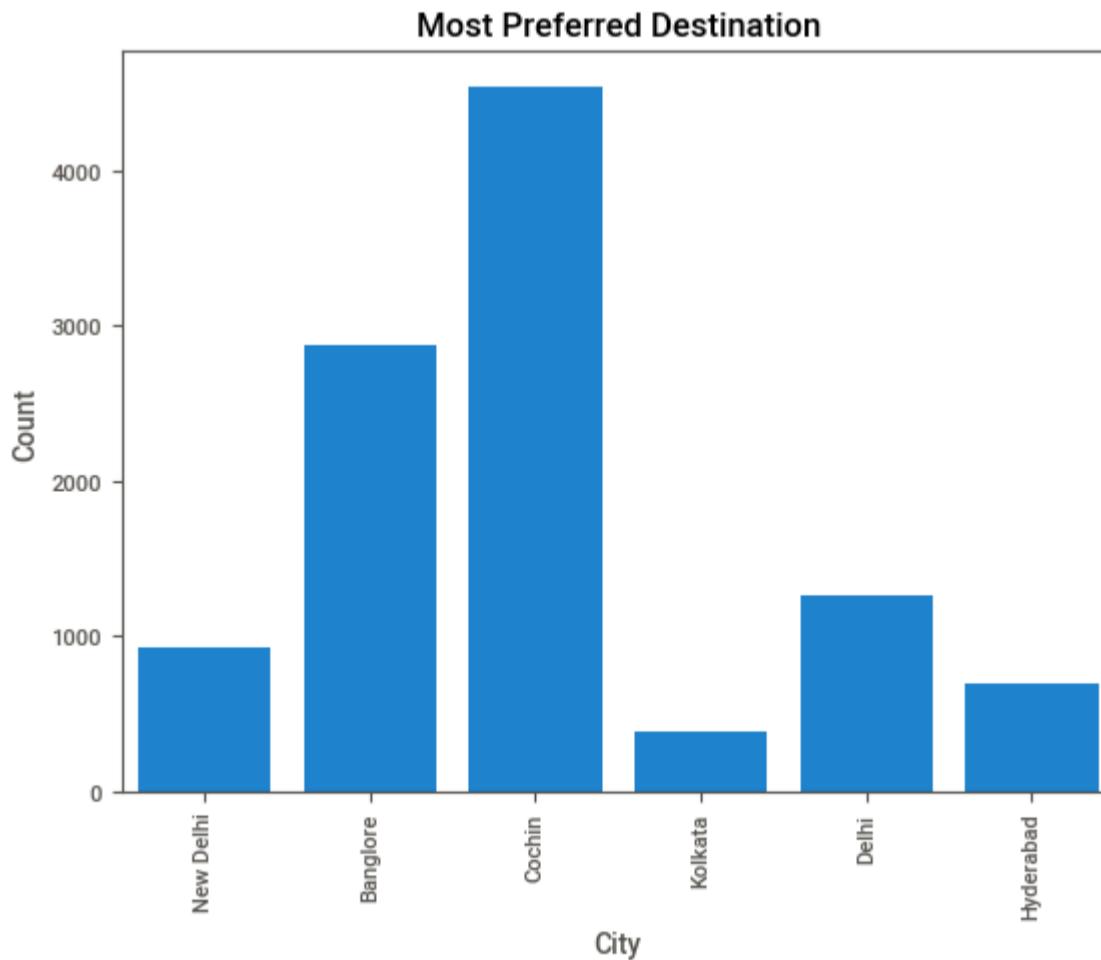
- From the Plot, it can be seen that the Most desired Source (Departure station), is Delhi with a count of almost 4500 passengers.
- And the least preferred Source station is Chennai with a count of nearly 300 passengers.

Most Preferred Destination

```
In [13]: sns.countplot(x = 'Destination', data = df)

plt.title('Most Preferred Destination', fontsize = 12, color = 'black')
plt.xlabel('City', fontsize = 10)
plt.ylabel('Count', fontsize = 10)
plt.xticks(rotation = 90)

plt.show()
```



Observations:

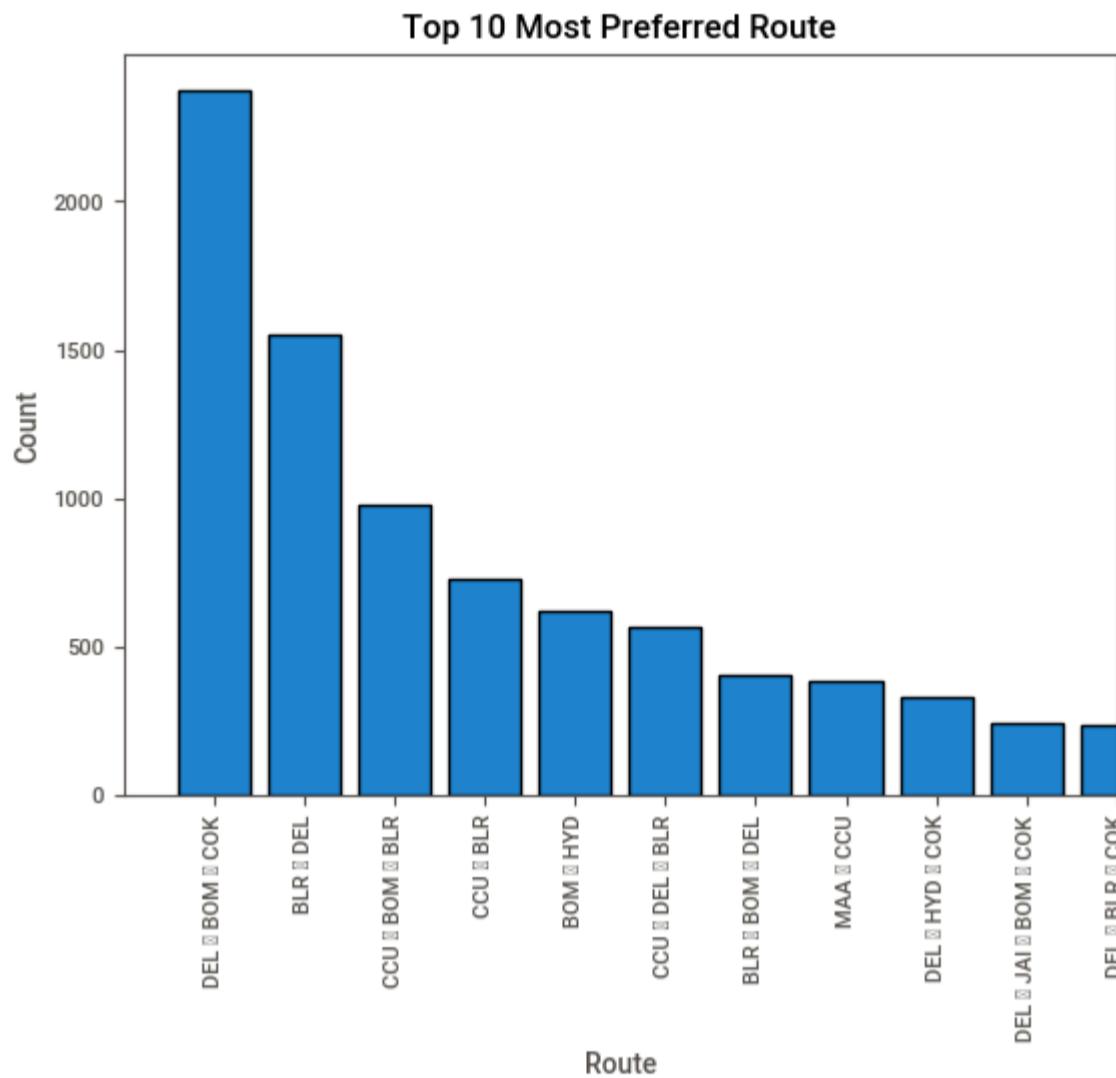
- From the plot it can be inferred that the Most Preferred Destination is Cochin city with a count of almost 4500 passengers.
- And Kolkata is the least preferred city among all with nearly 400 passengers.

TOP 10 MOST PREFERRED ROUTE FOR TRAVEL

```
In [14]: sns.countplot(x = 'Route', data = df, order = df['Route'].value_counts().index, ec = 'black')

plt.title('Top 10 Most Preferred Route', fontsize = 12, color = 'black')
plt.xlabel('Route', fontsize = 10)
plt.ylabel('Count', fontsize = 10)
plt.xticks(rotation = 90)
plt.xlim(-1, 10)

plt.show()
```



Observation:

- Most Preferred Route is DEL → BOM → COK with a count of nearly 3000 followed by BLR → DEL

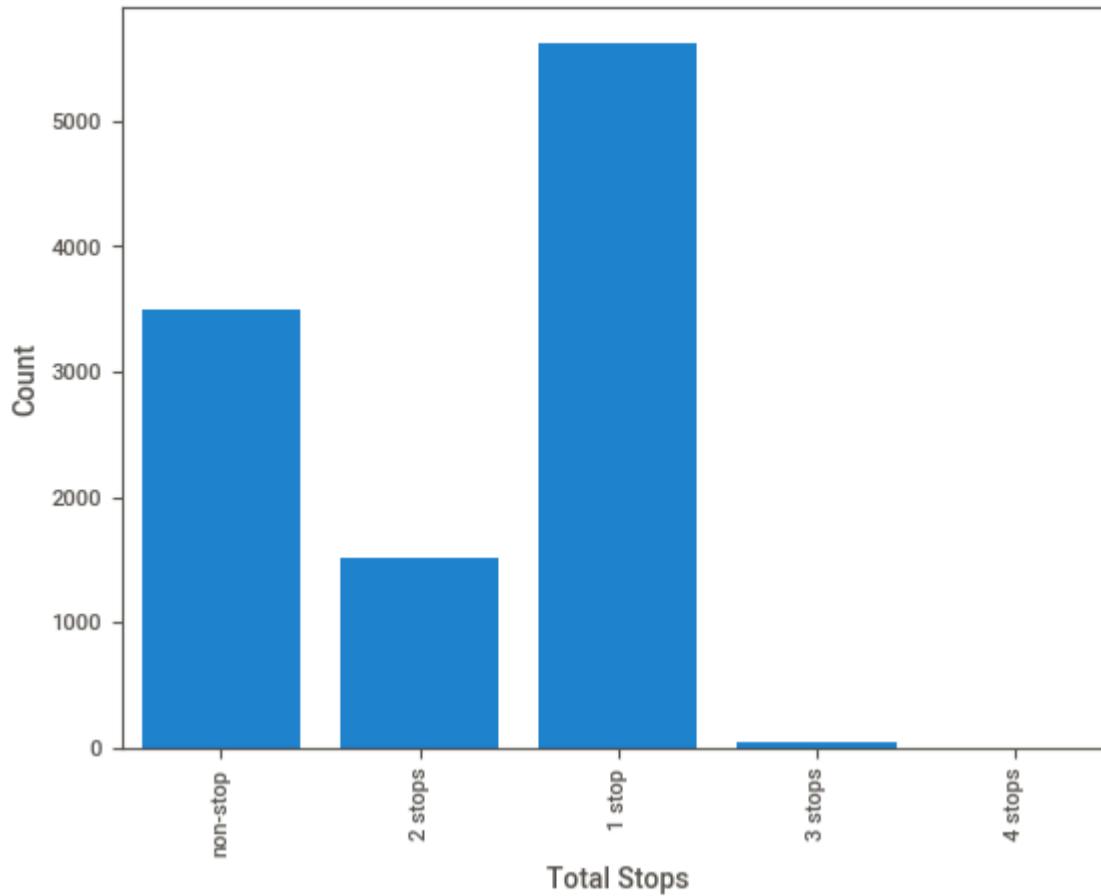
MOST NUMBER OF TOTAL_STOPS PREFERRED

```
In [15]: sns.countplot(x = 'Total_Stops', data = df)

plt.title('Most Total_Stops Preferred Route', fontsize = 12, color = 'black')
plt.xlabel('Total Stops', fontsize = 10)
plt.ylabel('Count', fontsize = 10)
plt.xticks(rotation = 90)

plt.show()
```

Most Total_Stops Preferred Route



Observations:

- From the plot, we can see that passengers preferred airlines route with 1 stop most with a count of almost 5800, followed by non-stop airline route.
- Passengers avoid 3 stops or 4 stops airline route.

Most Travelled Date

```
In [16]: plt.figure(figsize = (15, 10))

# Create a countplot to visualize the frequency of flights on different dates
sns.countplot(x = 'Date_of_Journey', data = df, order = df['Date_of_Journey'].value_counts().index)

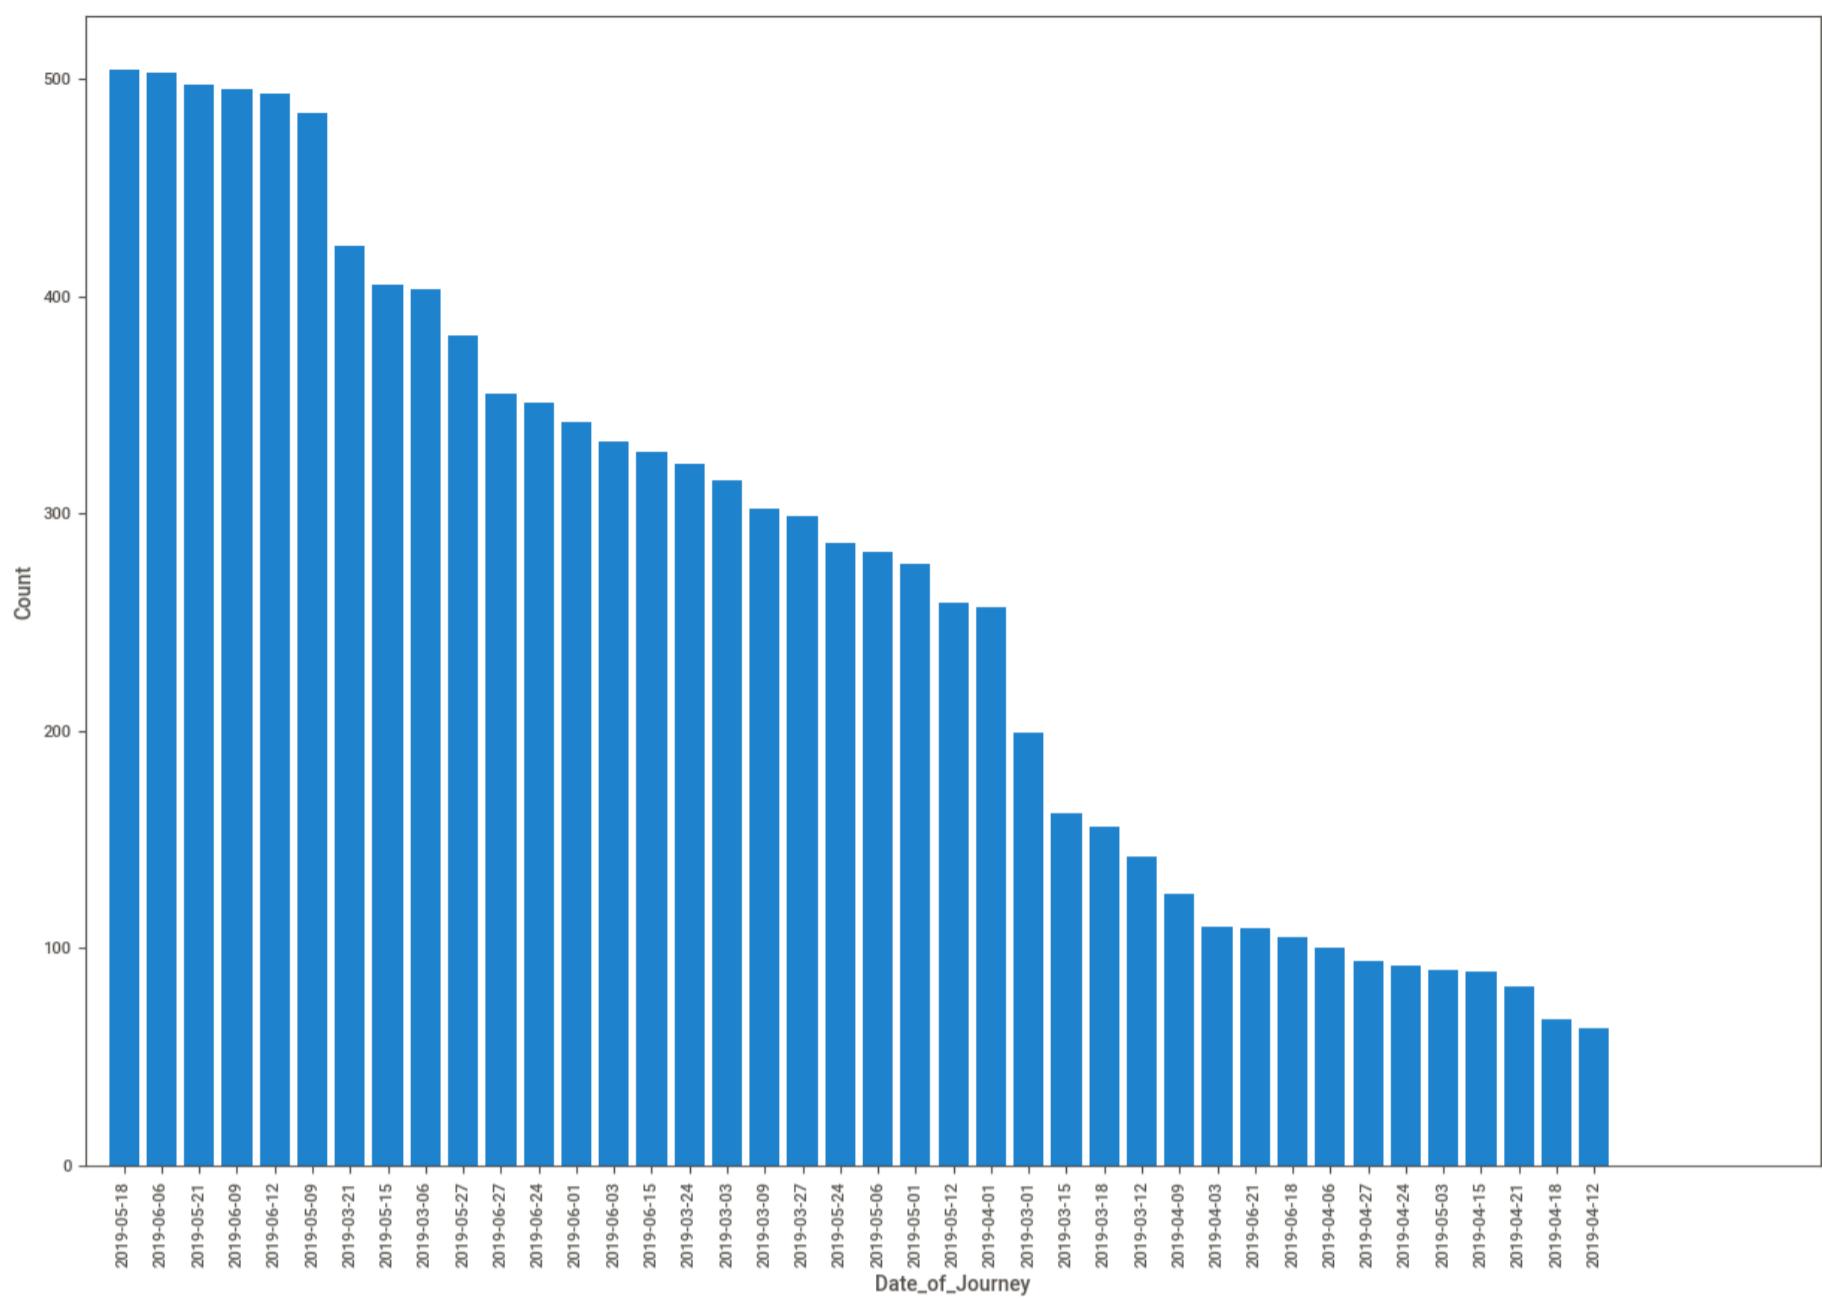
# Set x-axis and y-axis labels
plt.xlabel('Date_of_Journey', fontsize = 10)
plt.ylabel('Count', fontsize = 10)

# Rotate x-axis labels for better readability
plt.xticks(rotation = 90)

# Show the plot
plt.gca().invert_xaxis() # Invert the x-axis to display dates in descending order

# Limit the number of displayed dates on the x-axis if needed
plt.xlim(-1, 45)

# Show the plot
plt.show()
```



Observations:

- The number of passengers are higher on 18-05-2025 followed by 06-06-2025.
- The least number of passengers travelled was on 12-04-2025.

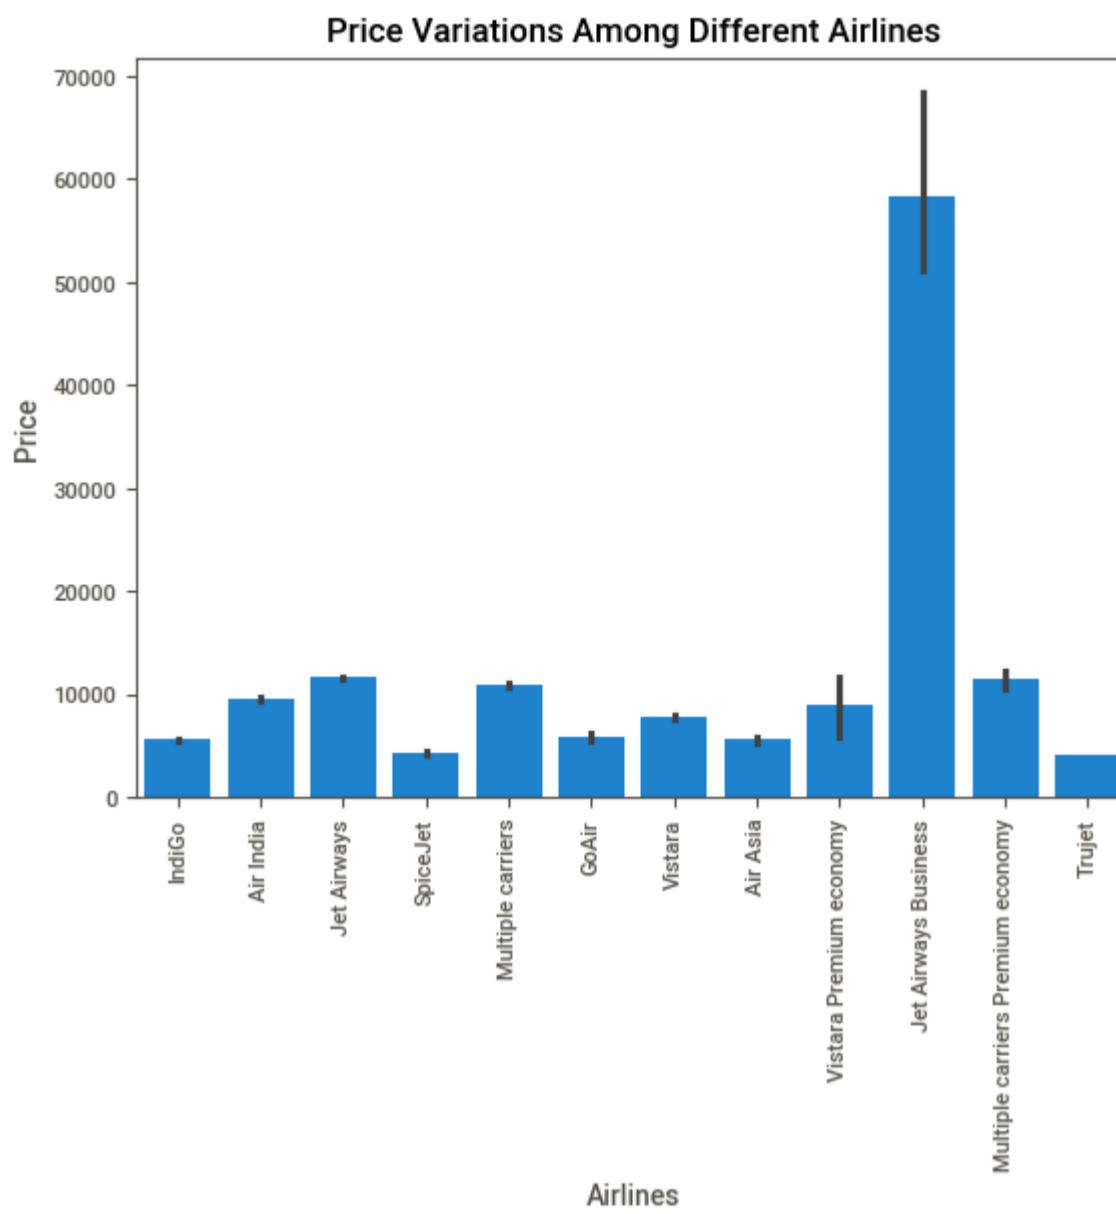
Bivariate Analysis

PRICE VARIATIONS BETWEEN AIRLINES

```
In [17]: sns.barplot(x = 'Airline', y = 'Price', data = df)

plt.title('Price Variations Among Different Airlines', fontsize = 12, color = 'black')
plt.xlabel('Airlines', fontsize = 10)
plt.ylabel('Price', fontsize = 10)
plt.xticks(rotation = 90)
```

```
Out[17]: ([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11],
[Text(0, 0, 'IndiGo'),
Text(1, 0, 'Air India'),
Text(2, 0, 'Jet Airways'),
Text(3, 0, 'SpiceJet'),
Text(4, 0, 'Multiple carriers'),
Text(5, 0, 'GoAir'),
Text(6, 0, 'Vistara'),
Text(7, 0, 'Air Asia'),
Text(8, 0, 'Vistara Premium economy'),
Text(9, 0, 'Jet Airways Business'),
Text(10, 0, 'Multiple carriers Premium economy'),
Text(11, 0, 'Trujet')])
```



Observations:

- Based on the plot, we can see that Jet Airways Business is the most expensive of all. Also, it is the least popular airline, as we observed in Airline Count Plot. This may be because of high price of the ticket.
- Jet Airways is the 2nd most expensive.
- Trujet is the least expensive airline.

RELATION BETWEEN PRICE AND TOTAL NUMBER OF STOPS

```
In [18]: sns.barplot(x = 'Total_Stops', y = 'Price', data = df)

plt.title('Relation Between Price and Total_Stops', fontsize = 12, color = 'black')
plt.xlabel('Total_Stops', fontsize = 10)
plt.ylabel('Price', fontsize = 10)
plt.xticks(rotation = 90)

plt.show()
```



Observations:

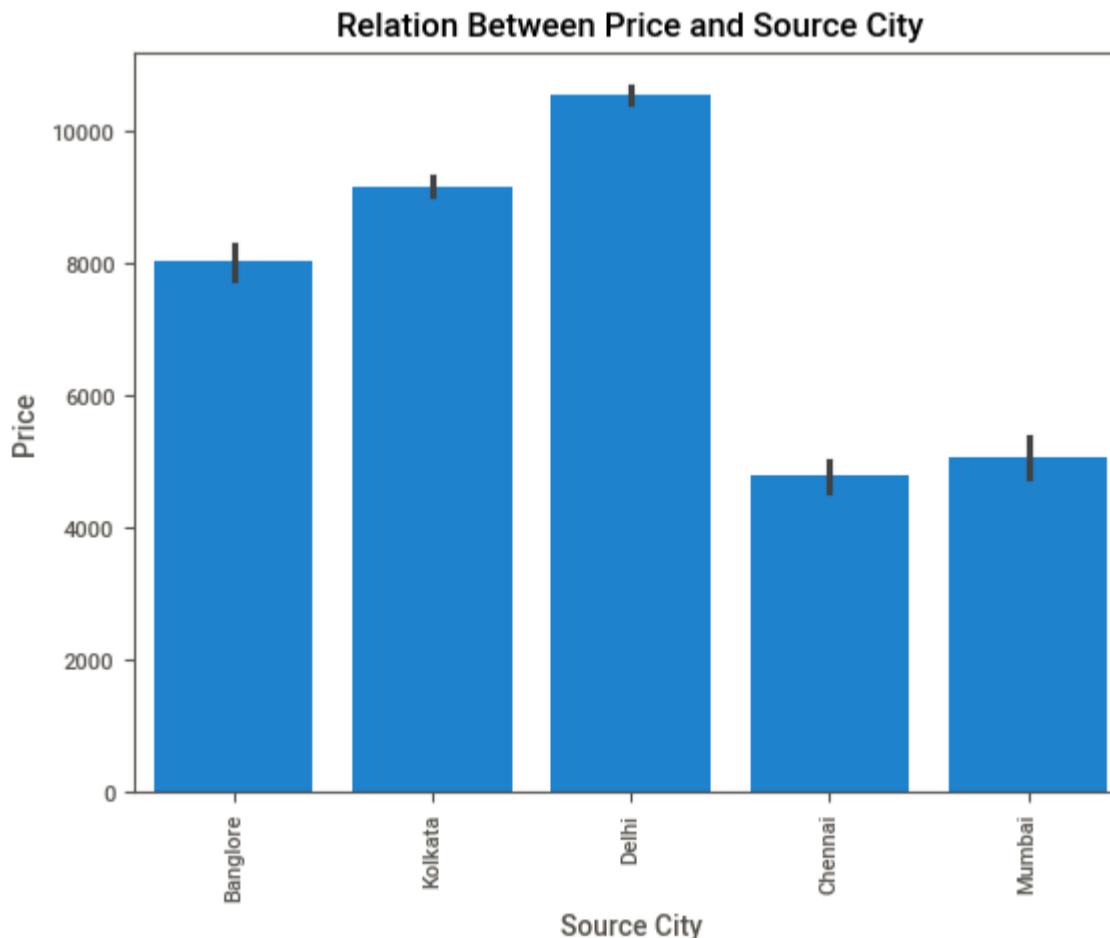
- The Price of Airline tickets with most number of stops are highest with rate nearly 17500.
- The least price is for the Airlines with no stops.

RELATION BETWEEN PRICE OF TICKETS AND THE SOURCE CITY

```
In [19]: sns.barplot(x = 'Source', y = 'Price', data = df)

plt.title('Relation Between Price and Source City', fontsize = 12, color = 'black')
plt.xlabel('Source City', fontsize = 10)
plt.ylabel('Price', fontsize = 10)
plt.xticks(rotation = 90)

plt.show()
```



Observations:

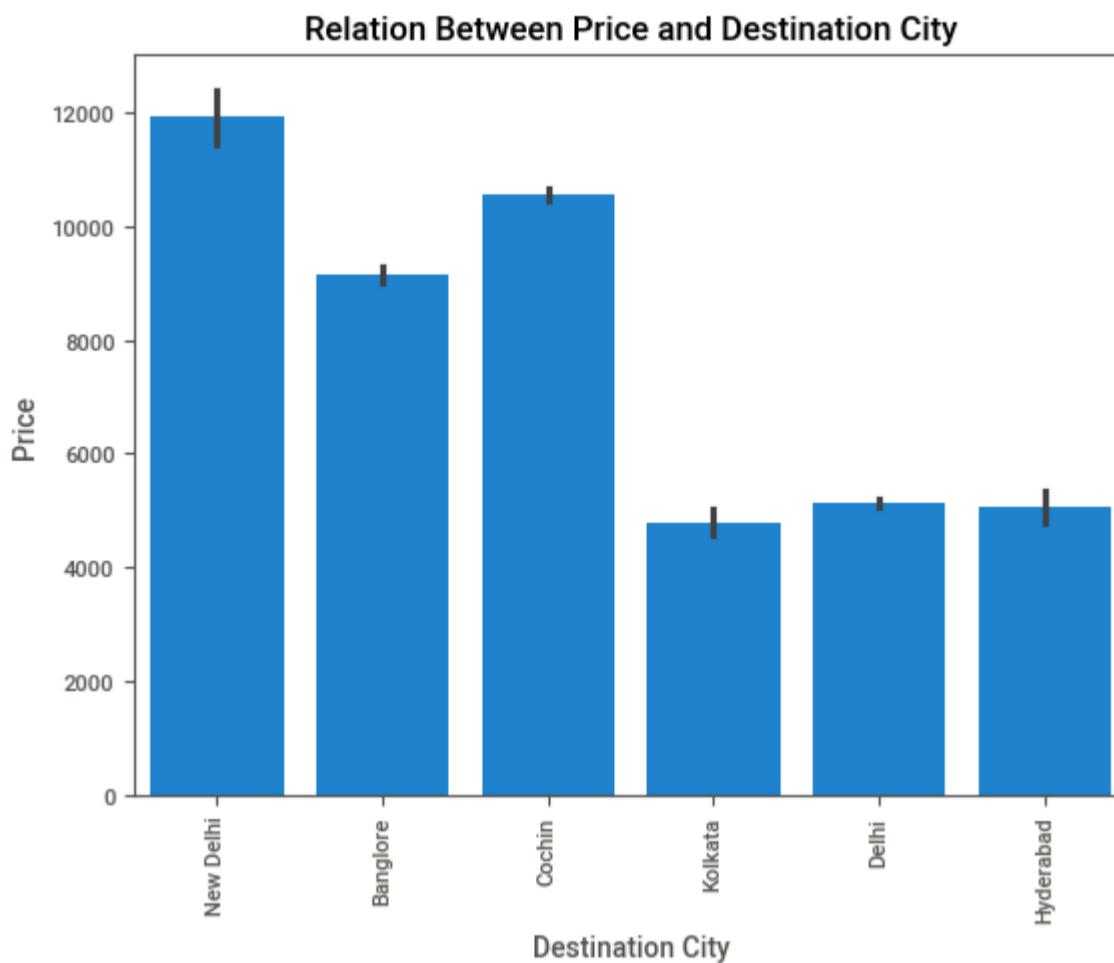
- The Price of Airline tickets is highest if the flight start from Delhi. Price of tickets is as high as 11000.
- The least price of ticket is from Chennai as source city. Price is as low as 4500.

RELATION BETWEEN PRICE OF TICKETS AND THE DESTINATION CITY

```
In [20]: sns.barplot(x = 'Destination', y = 'Price', data = df)

plt.title('Relation Between Price and Destination City', fontsize = 12, color = 'black')
plt.xlabel('Destination City', fontsize = 10)
plt.ylabel('Price', fontsize = 10)
plt.xticks(rotation = 90)

plt.show()
```



Observations:

- The Price of Airline tickets is highest if the flight destination is New Delhi. Price of tickets is as high as 12000.
- The least price of ticket is to Kolkata as destination city. Price is as low as 5000.

RELATION BETWEEN PRICE OF TICKETS AND THE ADDITIONAL INFORMATION PROVIDED

```
In [21]: # Listing the Unique values

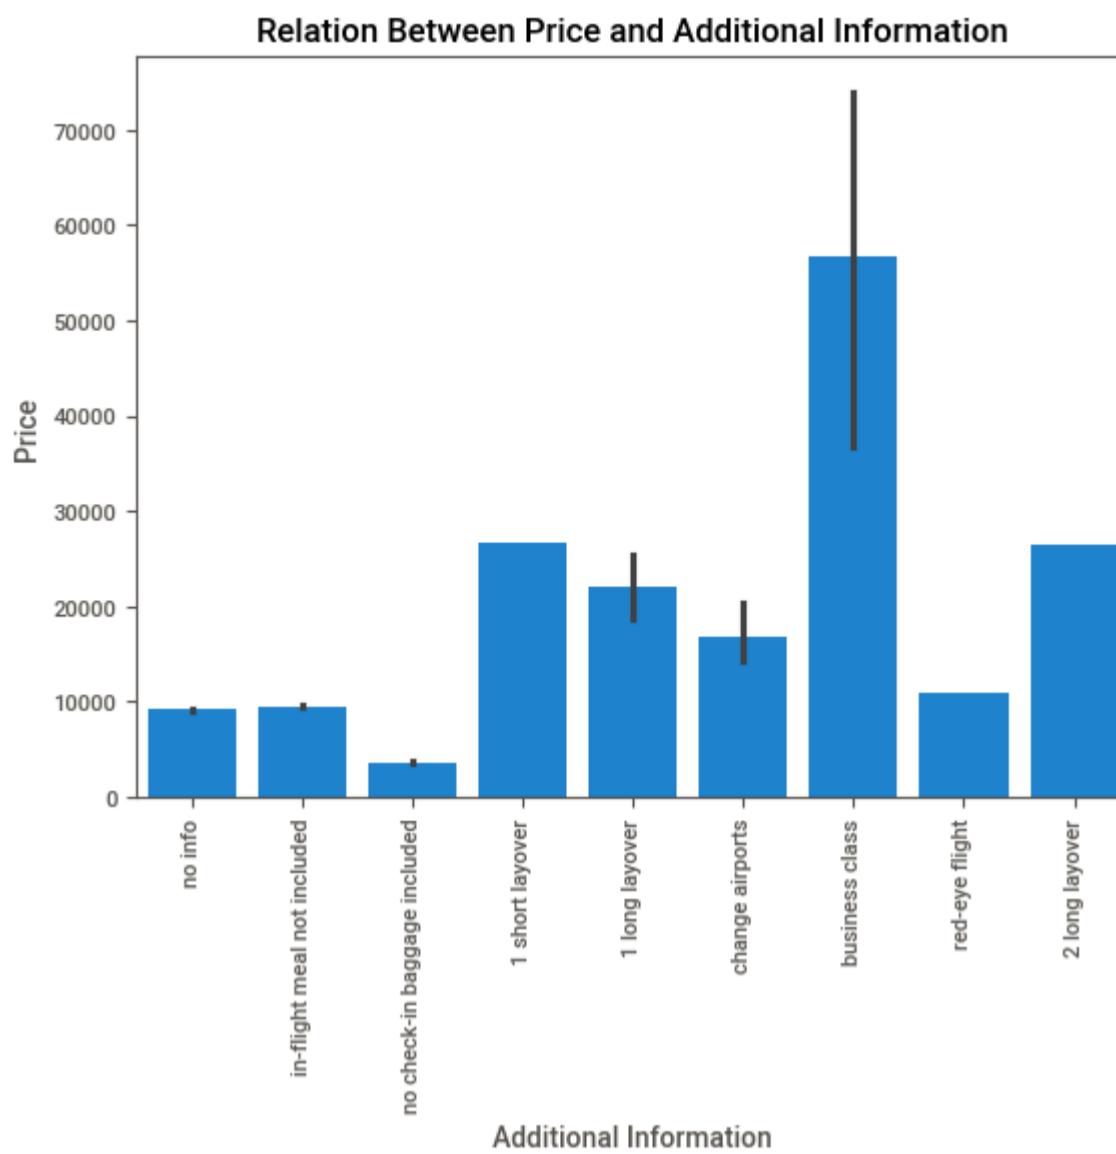
df.Additional_Info = list(df.Additional_Info.str.lower())
df.Additional_Info.unique()
```

```
Out[21]: array(['no info', 'in-flight meal not included',
       'no check-in baggage included', '1 short layover',
       '1 long layover', 'change airports', 'business class',
       'red-eye flight', '2 long layover'], dtype=object)
```

```
In [22]: sns.barplot(x = 'Additional_Info', y = 'Price', data = df)

plt.title('Relation Between Price and Additional Information', fontsize = 12, color = 'black')
plt.xlabel('Additional Information', fontsize = 10)
plt.ylabel('Price', fontsize = 10)
plt.xticks(rotation = 90)

plt.show()
```



Observations:

- The Price of Airline tickets is highest of Business Class category. This is followed by 1 Short Layover and 2 Short Layover.
- the Price is least for No Check-in Baggage Included.

COMPARISON OF PRICES ON WEEKDAYS AND ON WEEKENDS

```
In [23]: days_df = df[['Airline', 'Date_of_Journey', 'Price']]
days_df
```

```
Out[23]:
```

	Airline	Date_of_Journey	Price
0	IndiGo	2019-03-24	3897
1	Air India	2019-05-01	7662
2	Jet Airways	2019-06-09	13882
3	IndiGo	2019-05-12	6218
4	IndiGo	2019-03-01	13302
...
10678	Air Asia	2019-04-09	4107
10679	Air India	2019-04-27	4145
10680	Jet Airways	2019-04-27	7229
10681	Vistara	2019-03-01	12648
10682	Air India	2019-05-09	11753

10683 rows × 3 columns

```
In [24]: days_df['Date_of_Journey'] = pd.to_datetime(days_df['Date_of_Journey'], format = '%d/%m/%Y')
days_df['Weekday'] = days_df['Date_of_Journey'].dt.day_name()
days_df['Weekend'] = days_df['Weekday'].apply(lambda day:1 if day == 'Sunday' else 0)
days_df
```

Out[24]:

	Airline	Date_of_Journey	Price	Weekday	Weekend
0	IndiGo	2019-03-24	3897	Sunday	1
1	Air India	2019-05-01	7662	Wednesday	0
2	Jet Airways	2019-06-09	13882	Sunday	1
3	IndiGo	2019-05-12	6218	Sunday	1
4	IndiGo	2019-03-01	13302	Friday	0
...
10678	Air Asia	2019-04-09	4107	Tuesday	0
10679	Air India	2019-04-27	4145	Saturday	0
10680	Jet Airways	2019-04-27	7229	Saturday	0
10681	Vistara	2019-03-01	12648	Friday	0
10682	Air India	2019-05-09	11753	Thursday	0

10683 rows × 5 columns

In [25]:

```
# Creating a subplot with a suitable size
plt.subplots(figsize = (15, 10))

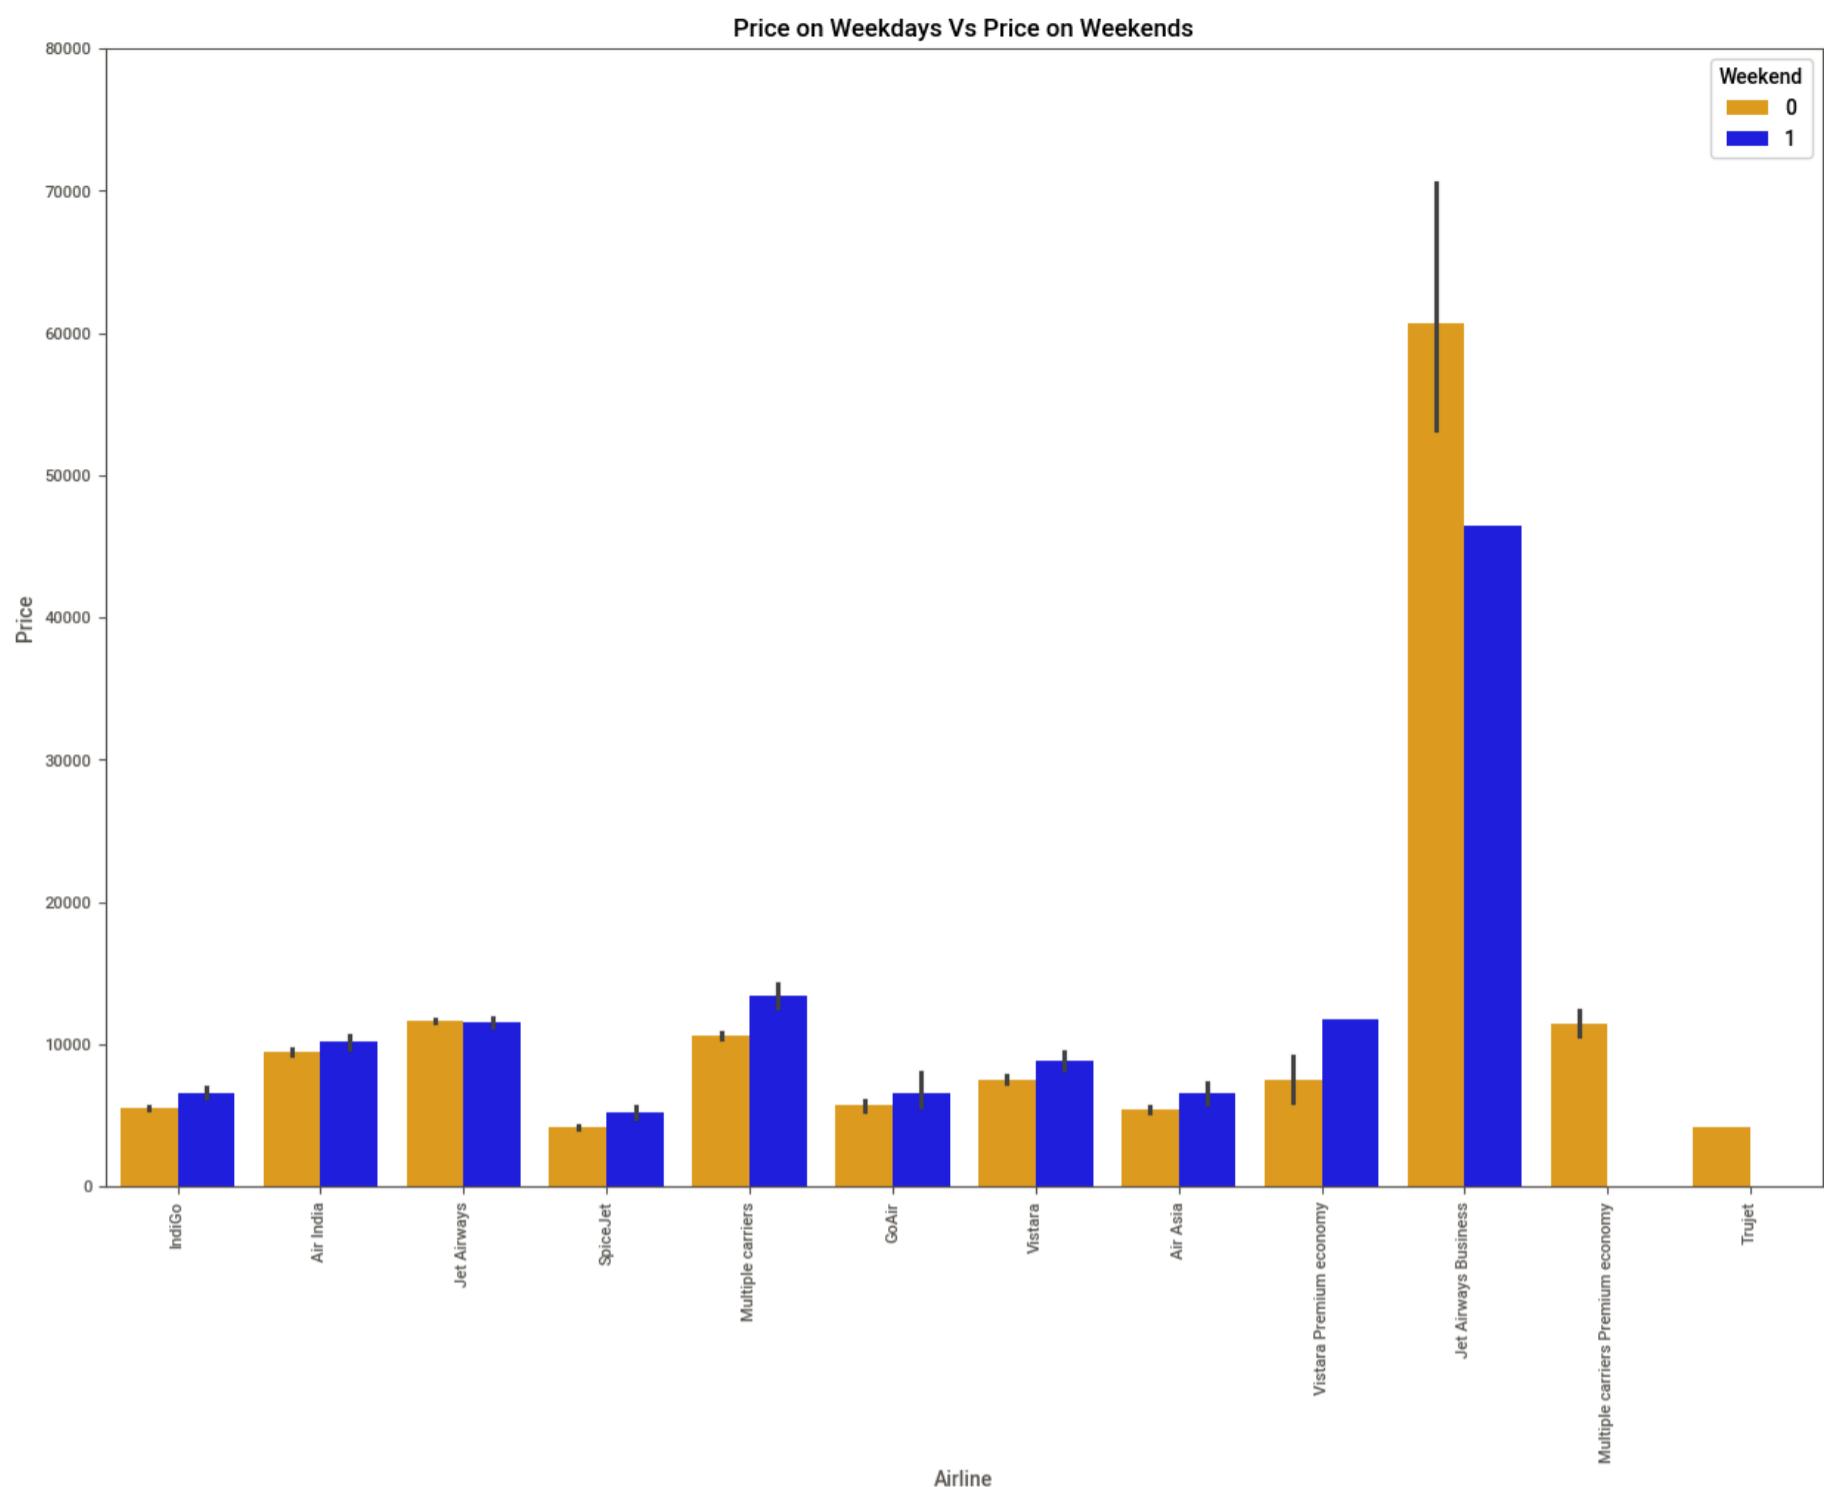
# Mapping the 'Weekend' column to colors (1 to 'blue' and 0 to 'orange')
color_palette = {1: 'blue', 0: 'orange'}

# Creating a barplot using Seaborn, mapping 'Weekend' to colors
sns.barplot(x = 'Airline', data = days_df, y = 'Price', hue = 'Weekend', palette = color_palette)

plt.title('Price on Weekdays Vs Price on Weekends', fontsize = 12, color = 'black')
plt.xlabel('Airline', fontsize = 10)
plt.ylabel('Price', fontsize = 10)
plt.xticks(rotation=90)

# Setting the y-axis limit
plt.ylim(0, 80000)

plt.show()
```



Observations:

- From the above plot, we can see that Jet Airways Business, Multiple Carriers Premium Economy have higher ticket prices during Weekdays than on Weekends.
- Other Airlines have higher ticket rates on Weekends than on Weekdays.
- Hence, it can be inferred that ticket prices for most of the Airlines are higher during Weekends than on the Weekdays.

Data Preprocessing And Feature Engineering

In [26]: df

Out[26]:

	Airline	Date_of_Journey	Source	Destination	Route	Dep_Time	Arrival_Time	Duration	Total_Stops	Additional_Info	Price
0	IndiGo	2019-03-24	Banglore	New Delhi	BLR → DEL	22:20	01:10 22 Mar	2h 50m	non-stop	no info	3897
1	Air India	2019-05-01	Kolkata	Banglore	CCU → IXR → BBI → BLR	05:50	13:15	7h 25m	2 stops	no info	7662
2	Jet Airways	2019-06-09	Delhi	Cochin	DEL → LKO → BOM → COK	09:25	04:25 10 Jun	19h	2 stops	no info	13882
3	IndiGo	2019-05-12	Kolkata	Banglore	CCU → NAG → BLR	18:05	23:30	5h 25m	1 stop	no info	6218
4	IndiGo	2019-03-01	Banglore	New Delhi	BLR → NAG → DEL	16:50	21:35	4h 45m	1 stop	no info	13302
...
10678	Air Asia	2019-04-09	Kolkata	Banglore	CCU → BLR	19:55	22:25	2h 30m	non-stop	no info	4107
10679	Air India	2019-04-27	Kolkata	Banglore	CCU → BLR	20:45	23:20	2h 35m	non-stop	no info	4145
10680	Jet Airways	2019-04-27	Banglore	Delhi	BLR → DEL	08:20	11:20	3h	non-stop	no info	7229
10681	Vistara	2019-03-01	Banglore	New Delhi	BLR → DEL	11:30	14:10	2h 40m	non-stop	no info	12648
10682	Air India	2019-05-09	Delhi	Cochin	DEL → GOI → BOM → COK	10:55	19:15	8h 20m	2 stops	no info	11753

10683 rows × 11 columns

Checking for Missing Values

In [27]: `df.isnull().sum()`

```
Out[27]: Airline      0
Date_of_Journey  0
Source          0
Destination     0
Route           1
Dep_Time        0
Arrival_Time    0
Duration        0
Total_Stops     1
Additional_Info 0
Price           0
dtype: int64
```

Observation:

- Since there is only one missing value in Route and Total_Stops columns, so these data will be dropped.

In [28]: `df.dropna(inplace = True)`

Rechecking for Missing Data after handling

```
In [29]: df.isnull().sum()
```

```
Out[29]: Airline      0
Date_of_Journey  0
Source        0
Destination    0
Route         0
Dep_Time       0
Arrival_Time   0
Duration       0
Total_Stops    0
Additional_Info 0
Price          0
dtype: int64
```

Observation:

- No Missing data present

Converting 'Date of Journey' Column to Numeric Data

```
In [30]: # Getting the Unique Dates
df.Date_of_Journey.unique()
```

```
Out[30]: <DatetimeArray>
['2019-03-24 00:00:00', '2019-05-01 00:00:00', '2019-06-09 00:00:00',
 '2019-05-12 00:00:00', '2019-03-01 00:00:00', '2019-06-24 00:00:00',
 '2019-03-12 00:00:00', '2019-05-27 00:00:00', '2019-06-01 00:00:00',
 '2019-04-18 00:00:00', '2019-05-09 00:00:00', '2019-04-24 00:00:00',
 '2019-03-03 00:00:00', '2019-04-15 00:00:00', '2019-06-12 00:00:00',
 '2019-03-06 00:00:00', '2019-03-21 00:00:00', '2019-04-03 00:00:00',
 '2019-05-06 00:00:00', '2019-05-15 00:00:00', '2019-06-18 00:00:00',
 '2019-06-15 00:00:00', '2019-04-06 00:00:00', '2019-05-18 00:00:00',
 '2019-06-27 00:00:00', '2019-05-21 00:00:00', '2019-06-03 00:00:00',
 '2019-03-15 00:00:00', '2019-05-03 00:00:00', '2019-03-09 00:00:00',
 '2019-06-06 00:00:00', '2019-05-24 00:00:00', '2019-04-01 00:00:00',
 '2019-04-21 00:00:00', '2019-06-21 00:00:00', '2019-03-27 00:00:00',
 '2019-03-18 00:00:00', '2019-04-12 00:00:00', '2019-04-09 00:00:00',
 '2019-04-27 00:00:00']
Length: 40, dtype: datetime64[ns]
```

```
In [31]: # Dropping the Date_of_Journey Column
df = df.drop(['Date_of_Journey'], axis = 1)
```

```
In [32]: df.head(2)
```

```
Out[32]:   Airline  Source  Destination      Route  Dep_Time  Arrival_Time  Duration  Total_Stops  Additional_Info  Price
0  IndiGo  Banglore     New Delhi  BLR → DEL  22:20  01:10 22 Mar  2h 50m  non-stop      no info  3897
1  Air India  Kolkata     Banglore  CCU → IXR → BBI → BLR  05:50  13:15  7h 25m  2 stops      no info  7662
```

Converting 'Dep_Time' Column to Numeric Data

```
In [33]: # Splitting the Dep_Time into two different columns - one containing Hours and other containing Minutes
df["Dep_Hour"] = pd.to_datetime(df["Dep_Time"]).dt.hour
df["Dep_Min"] = pd.to_datetime(df["Dep_Time"]).dt.minute
```

```
In [34]: # Dropping the Dep_Time Column
df = df.drop(['Dep_Time'], axis = 1)
```

```
In [35]: df.head(2)
```

```
Out[35]:   Airline  Source  Destination      Route  Arrival_Time  Duration  Total_Stops  Additional_Info  Price  Dep_Hour  Dep_Min
0  IndiGo  Banglore     New Delhi  BLR → DEL  01:10 22 Mar  2h 50m  non-stop      no info  3897        22        20
1  Air India  Kolkata     Banglore  CCU → IXR → BBI → BLR  13:15  7h 25m  2 stops      no info  7662         5        50
```

Converting 'Arrival_Time' Column to Numeric Data

```
In [36]: # Splitting the Dep_Time into two different columns - one containing Hours and other containing Minutes
df["Arrival_Hour"] = pd.to_datetime(df["Arrival_Time"]).dt.hour
df["Arrival_Min"] = pd.to_datetime(df["Arrival_Time"]).dt.minute
```

```
In [37]: # Dropping the Arrival_Time Column
```

```
df = df.drop(['Arrival_Time'], axis = 1)
```

```
In [38]: df.head(2)
```

```
Out[38]:
```

	Airline	Source	Destination	Route	Duration	Total_Stops	Additional_Info	Price	Dep_Hour	Dep_Min	Arrival_Hour	Arrival_Min
0	IndiGo	Banglore	New Delhi	BLR → DEL	2h 50m	non-stop	no info	3897	22	20	1	10
1	Air India	Kolkata	Banglore	CCU → IXR → BBI → BLR	7h 25m	2 stops	no info	7662	5	50	13	15

Converting 'Duration' (hh:mm) format to 'Duration' (Minutes) Format

```
In [39]: df['Duration'] = df['Duration'].str.replace("h", '*60').str.replace(' ','+') .str.replace('m','*1').apply(eval)
```

```
In [40]: df.head(2)
```

```
Out[40]:
```

	Airline	Source	Destination	Route	Duration	Total_Stops	Additional_Info	Price	Dep_Hour	Dep_Min	Arrival_Hour	Arrival_Min
0	IndiGo	Banglore	New Delhi	BLR → DEL	170	non-stop	no info	3897	22	20	1	10
1	Air India	Kolkata	Banglore	CCU → IXR → BBI → BLR	445	2 stops	no info	7662	5	50	13	15

Correlation among Features

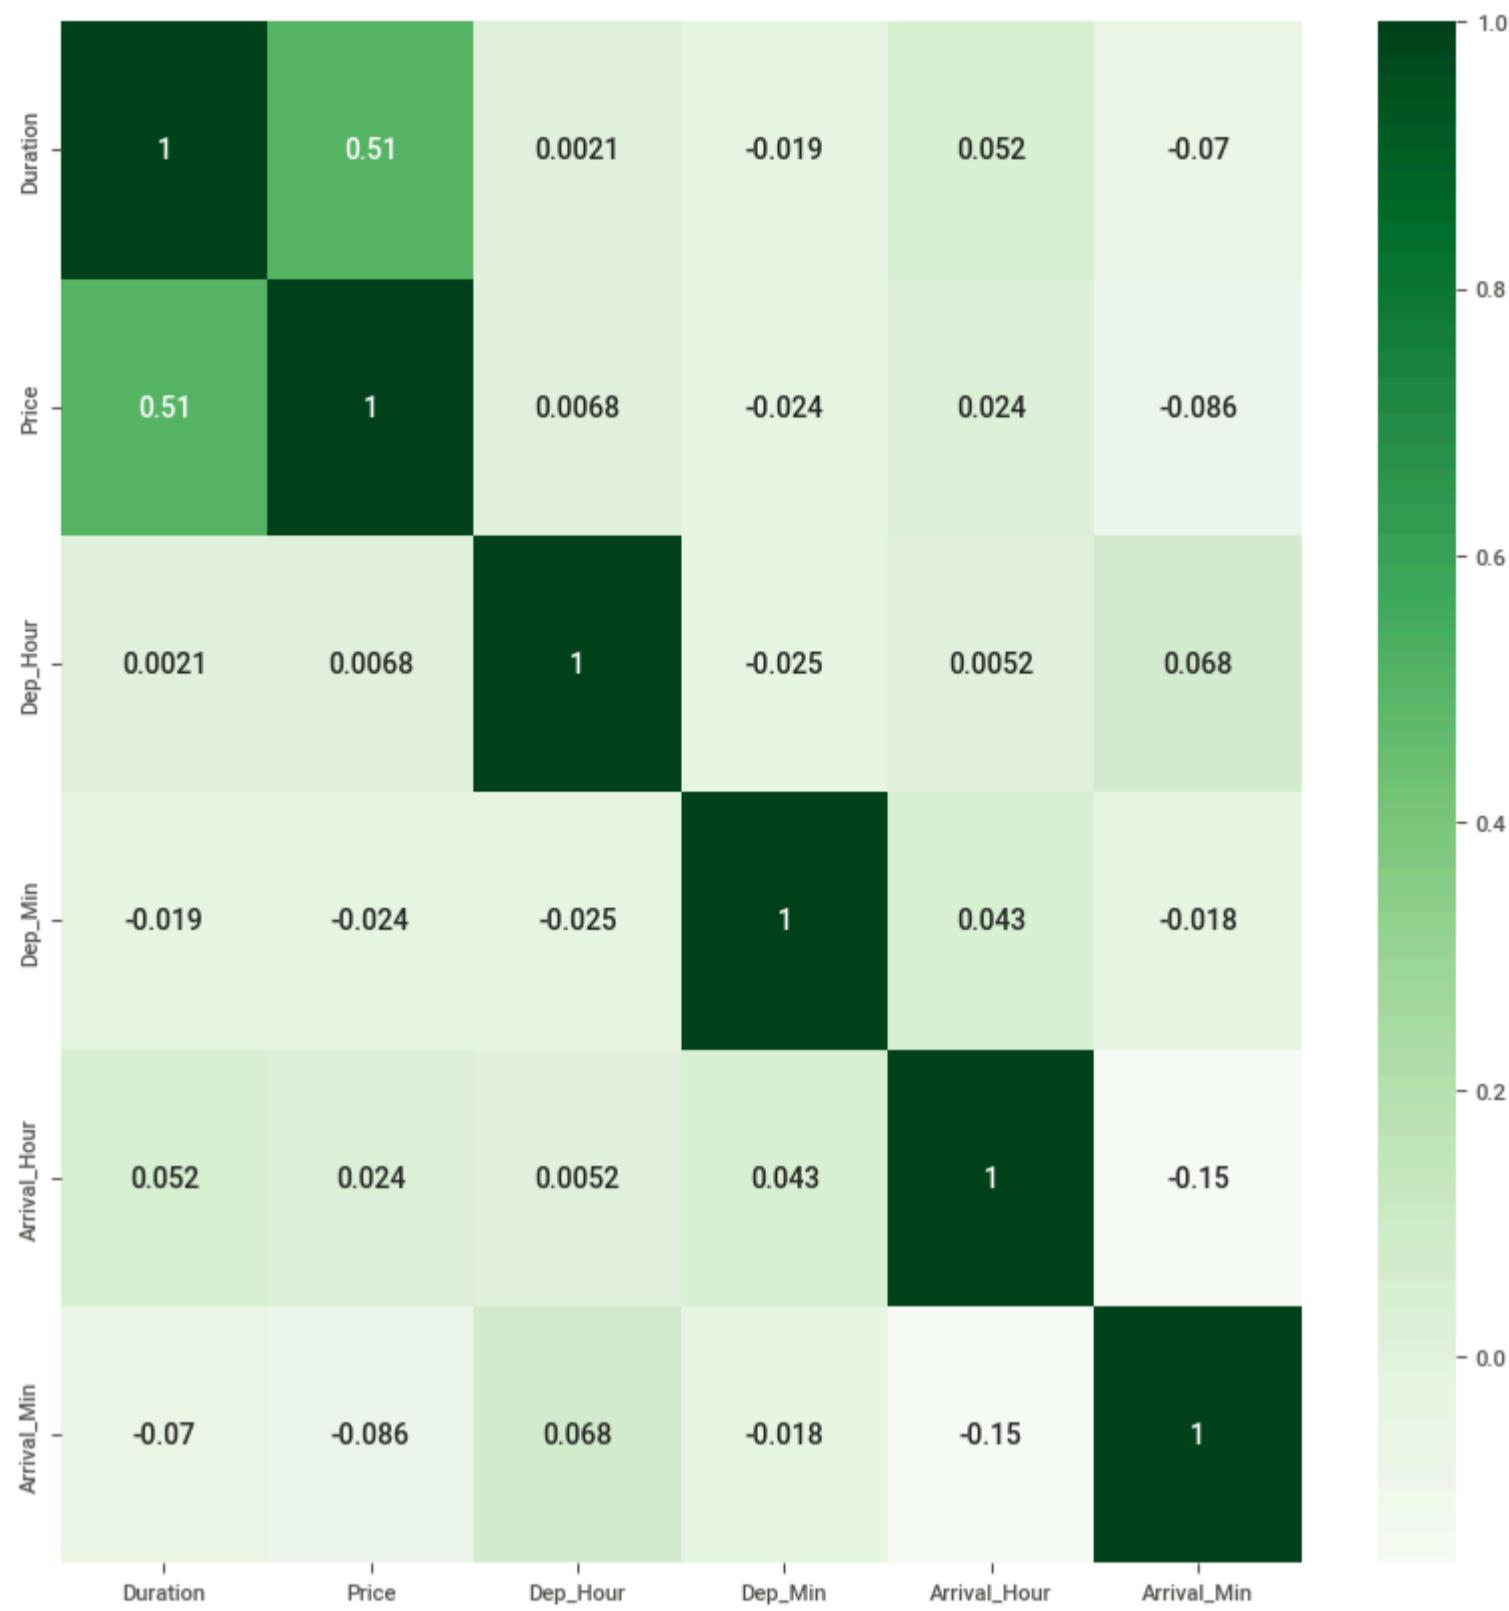
```
In [41]: df.corr(numeric_only=True)
```

```
Out[41]:
```

	Duration	Price	Dep_Hour	Dep_Min	Arrival_Hour	Arrival_Min
Duration	1.000000	0.506480	0.002088	-0.019099	0.051531	-0.069663
Price	0.506480	1.000000	0.006799	-0.024458	0.024244	-0.086155
Dep_Hour	0.002088	0.006799	1.000000	-0.024745	0.005180	0.067911
Dep_Min	-0.019099	-0.024458	-0.024745	1.000000	0.043122	-0.017597
Arrival_Hour	0.051531	0.024244	0.005180	0.043122	1.000000	-0.154363
Arrival_Min	-0.069663	-0.086155	0.067911	-0.017597	-0.154363	1.000000

Plotting Heatmap

```
In [42]: plt.figure(figsize = (10, 10))
sns.heatmap(df.corr(numeric_only=True), cmap = 'Greens', annot = True)
plt.show()
```



Observation:

- No significant correlation between features

Converting Categorical Data to Numerical Data

In [43]: df

Out[43]:

	Airline	Source	Destination	Route	Duration	Total_Stops	Additional_Info	Price	Dep_Hour	Dep_Min	Arrival_Hour	Arrival_M
0	IndiGo	Banglore	New Delhi	BLR → DEL	170	non-stop	no info	3897	22	20	1	
1	Air India	Kolkata	Banglore	CCU → IXR → BBI → BLR	445	2 stops	no info	7662	5	50	13	
2	Jet Airways	Delhi	Cochin	DEL → LKO → BOM → COK	1140	2 stops	no info	13882	9	25	4	
3	IndiGo	Kolkata	Banglore	CCU → NAG → BLR	325	1 stop	no info	6218	18	5	23	
4	IndiGo	Banglore	New Delhi	BLR → NAG → DEL	285	1 stop	no info	13302	16	50	21	
...	
10678	Air Asia	Kolkata	Banglore	CCU → BLR	150	non-stop	no info	4107	19	55	22	
10679	Air India	Kolkata	Banglore	CCU → BLR	155	non-stop	no info	4145	20	45	23	
10680	Jet Airways	Banglore	Delhi	BLR → DEL	180	non-stop	no info	7229	8	20	11	
10681	Vistara	Banglore	New Delhi	BLR → DEL	160	non-stop	no info	12648	11	30	14	
10682	Air India	Delhi	Cochin	DEL → GOI → BOM → COK	500	2 stops	no info	11753	10	55	19	

10682 rows × 12 columns

Observation:

- Following Columns are Categorical in nature:
 1. Airline
 2. Source
 3. Destination
 4. Route
 5. Total_stops
 6. Additional_Info
- These columns need to be handled and converted to numeric type

In [44]: `!pip install scikit-learn`

```
Requirement already satisfied: scikit-learn in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (1.7.2)
Requirement already satisfied: numpy>=1.22.0 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from scikit-learn) (2.3.3)
Requirement already satisfied: scipy>=1.8.0 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from scikit-learn) (1.16.2)
Requirement already satisfied: joblib>=1.2.0 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from scikit-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from scikit-learn) (3.6.0)
[notice] A new release of pip is available: 25.2 -> 25.3
[notice] To update, run: python.exe -m pip install --upgrade pip
```

Using Label Encoder to convert Categorical Data to Numeric Data

```
In [45]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()

# Creating a Function to encode Categorical Data

def encode(columns,data):

    for column in columns:
        data[column] = le.fit_transform(data[column])

    return data

# Passing the Categorical Columns to the 'encode' function

obj = df[['Airline', 'Source', 'Destination', 'Route', 'Total_Stops', 'Additional_Info']]

df = encode(obj, df)
```

```
In [46]: df
```

	Airline	Source	Destination	Route	Duration	Total_Stops	Additional_Info	Price	Dep_Hour	Dep_Min	Arrival_Hour	Arrival_Min
0	3	0	5	18	170	4	7	3897	22	20	1	10
1	1	3	0	84	445	1	7	7662	5	50	13	15
2	4	2	1	118	1140	1	7	13882	9	25	4	25
3	3	3	0	91	325	0	7	6218	18	5	23	30
4	3	0	5	29	285	0	7	13302	16	50	21	35
...
10678	0	3	0	64	150	4	7	4107	19	55	22	25
10679	1	3	0	64	155	4	7	4145	20	45	23	20
10680	4	0	2	18	180	4	7	7229	8	20	11	20
10681	10	0	5	18	160	4	7	12648	11	30	14	10
10682	1	2	1	108	500	1	7	11753	10	55	19	15

10682 rows × 12 columns

Scaling Data

```
In [47]: x = df.drop('Price', axis = 1)
x
```

Out[47]:

	Airline	Source	Destination	Route	Duration	Total_Stops	Additional_Info	Dep_Hour	Dep_Min	Arrival_Hour	Arrival_Min
0	3	0	5	18	170	4	7	22	20	1	10
1	1	3	0	84	445	1	7	5	50	13	15
2	4	2	1	118	1140	1	7	9	25	4	25
3	3	3	0	91	325	0	7	18	5	23	30
4	3	0	5	29	285	0	7	16	50	21	35
...
10678	0	3	0	64	150	4	7	19	55	22	25
10679	1	3	0	64	155	4	7	20	45	23	20
10680	4	0	2	18	180	4	7	8	20	11	20
10681	10	0	5	18	160	4	7	11	30	14	10
10682	1	2	1	108	500	1	7	10	55	19	15

10682 rows × 11 columns

In [48]:

```
y = df.Price
y
```

Out[48]:

```
0      3897
1      7662
2     13882
3      6218
4     13302
...
10678    4107
10679    4145
10680    7229
10681   12648
10682   11753
Name: Price, Length: 10682, dtype: int64
```

MinMax Scalar

In [49]:

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
x_scaled = scaler.fit_transform(x)
```

In [50]:

```
x_scaled
```

Out[50]:

```
array([[0.27272727, 0.          , 1.          , ... , 0.36363636, 0.04347826,
       0.18181818],
       [0.09090909, 0.75        , 0.          , ... , 0.90909091, 0.56521739,
       0.27272727],
       [0.36363636, 0.5        , 0.2        , ... , 0.45454545, 0.17391304,
       0.45454545],
       ... ,
       [0.36363636, 0.          , 0.4        , ... , 0.36363636, 0.47826087,
       0.36363636],
       [0.90909091, 0.          , 1.          , ... , 0.54545455, 0.60869565,
       0.18181818],
       [0.09090909, 0.5        , 0.2        , ... , 1.          , 0.82608696,
       0.27272727]], shape=(10682, 11))
```

BUSINESS CASE 2: To Create a Predictive Model which will help the Customers to Predict the Future Flight Prices and Plan their Journey accordingly.

Model Creation

Creating Training Data and Testing Data Set

In [51]:

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size = 0.25, random_state = 42)
```

In [52]:

```
x_train.shape
```

Out[52]:

```
(8011, 11)
```

In [53]:

```
y_train.shape
```

```
Out[53]: (8011,)
```

```
In [54]: x_test.shape
```

```
Out[54]: (2671, 11)
```

```
In [55]: y_test.shape
```

```
Out[55]: (2671,)
```

Defining a Function for Model Building and its Evaluation

```
In [56]: from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

def evaluation_score(model):
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)

    mse = mean_squared_error(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    adj_r = 1 - (1 - r2) * (len(y_test) - 1) / (len(y_test) - x_test.shape[1] - 1)

    evaluation = {
        ('The Mean Squared Error for this algorithm is :', mse),
        ('The Mean Absolute Error for this algorithm is :', mae),
        ('The R2 score for this algorithm is          :', r2),
        ('The Adjusted R2 score for this algorithm is:', adj_r)
    }

    return evaluation
```

1. LINEAR REGRESSION ALGORITHM

```
In [57]: from sklearn.linear_model import LinearRegression

lr = LinearRegression()
```

```
In [58]: evaluation_score(lr)
```

```
Out[58]: {('The Adjusted R2 score for this algorithm is      :', 0.4150945849766562),
          ('The Mean Absolute Error for this algorithm is :', 2490.2569480836055),
          ('The Mean Squared Error for this algorithm is   :', 12006999.094906157),
          ('The R2 score for this algorithm is           :', 0.417504307660273)}
```

2. SUPPORT VECTOR MACHINE ALGORITHM

```
In [59]: from sklearn.svm import SVC
```

```
In [60]: evaluation_score(SVC())
```

```
Out[60]: {('The Adjusted R2 score for this algorithm is      :', 0.39993638667635356),
          ('The Mean Absolute Error for this algorithm is :', 1938.6173717708723),
          ('The Mean Squared Error for this algorithm is   :', 12318168.163609136),
          ('The R2 score for this algorithm is           :', 0.4024085588660764)}
```

3. K-NEAREST NEIGHBOURS ALGORITHM

```
In [61]: from sklearn.neighbors import KNeighborsRegressor
```

```
In [62]: ERROR_RATE = []

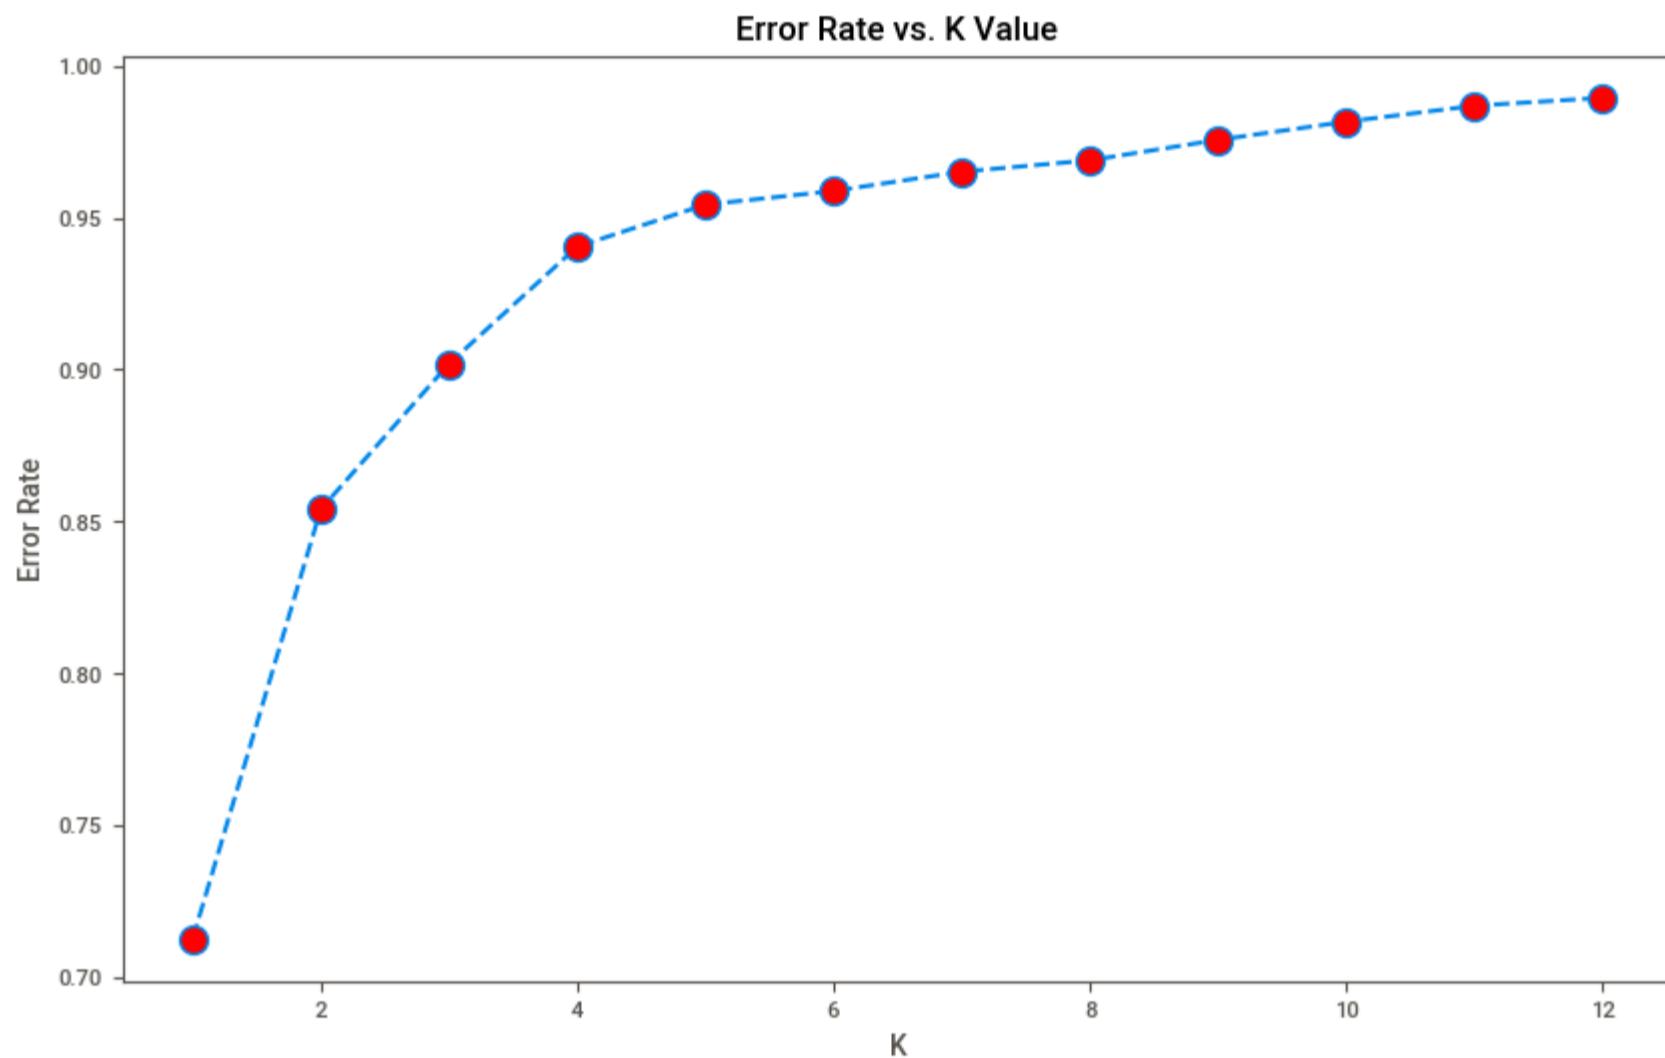
for i in range(1, 13):
    KNN = KNeighborsRegressor(i)
    KNN.fit(x_train, y_train)
    y_pred = KNN.predict(x_test)
    error_rate = (y_test != y_pred).sum() / len(y_test)
    ERROR_RATE.append(error_rate)
```

```
In [63]: ERROR_RATE = [float(err) for err in ERROR_RATE]
```

```
In [64]: ERROR_RATE
```

```
Out[64]: [0.7124672407338075,
 0.8539872706851367,
 0.901909397229502,
 0.9404717334331711,
 0.9543242231374017,
 0.958816922500936,
 0.9651815799326096,
 0.968925496068888,
 0.9756645451141894,
 0.9816548109322352,
 0.9868962935230251,
 0.98951703481842]
```

```
In [65]: plt.figure(figsize = (10, 6))
plt.plot(range(1, 13), ERROR_RATE, marker = 'o', linestyle = 'dashed', markerfacecolor = 'red', markersize = 10)
plt.title('Error Rate vs. K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')
plt.show()
```



```
In [66]: KNN = KNeighborsRegressor(n_neighbors = 5)
```

```
In [67]: evaluation_score(KNN)
```

```
Out[67]: {('The Adjusted R2 score for this alogorithm is   :', 0.7260703827491823),
 ('The Mean Absolute Error for this alogorithm is   :', 1338.8279296143767),
 ('The Mean Squared Error for this alogorithm is   :', 5623255.6271808315),
 ('The R2 score for this alogorithm is           :', 0.7271989317341108)}
```

4. DECISION TREE ALGORITHM

```
In [68]: from sklearn.tree import DecisionTreeRegressor
dtr = DecisionTreeRegressor(max_depth = 10, criterion = 'absolute_error')
```

```
In [69]: evaluation_score(dtr)
```

```
Out[69]: {('The Adjusted R2 score for this alogorithm is   :', 0.6423069112636544),
 ('The Mean Absolute Error for this alogorithm is   :', 1326.105578435043),
 ('The Mean Squared Error for this alogorithm is   :', 7342760.867652564),
 ('The R2 score for this alogorithm is           :', 0.6437805532022685)}
```

Hyperparameter Tuning for Decision Tree Algorithm

```
In [70]: from sklearn.model_selection import GridSearchCV

param_grid = {
    'max_depth': [None, 5, 10, 15],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt', 'log2']
}
```

```
grid_dtr = GridSearchCV(estimator = dtr, param_grid = param_grid, cv = 5, scoring = 'neg_mean_squared_error', n_jobs = -1)
grid_dtr.fit(x_train, y_train)

# Best Hyperparameters

best_params = grid_dtr.best_params_
print(f"Best Hyperparameters: {best_params}")

# Best Model

best_dtr = grid_dtr.best_estimator_

y_pred = best_dtr.predict(x_test)
```

Best Hyperparameters: {'max_depth': 15, 'max_features': 'log2', 'min_samples_leaf': 2, 'min_samples_split': 10}

In [71]: `evaluation_score(grid_dtr)`

Out[71]: {('The Adjusted R2 score for this alogorithm is :', 0.5897847875754627),
 ('The Mean Absolute Error for this alogorithm is :', 1361.9807188318982),
 ('The Mean Squared Error for this alogorithm is :', 8420940.476506926),
 ('The R2 score for this alogorithm is :', 0.5914748127951892)}

5. RANDOM FOREST ALGORITHM

In [72]: `from sklearn.ensemble import RandomForestRegressor`

In [73]: `rfr = RandomForestRegressor()`

In [74]: `evaluation_score(rfr)`

Out[74]: {('The Adjusted R2 score for this alogorithm is :', 0.6836355509866177),
 ('The Mean Absolute Error for this alogorithm is :', 1314.9071870862158),
 ('The Mean Squared Error for this alogorithm is :', 6494362.2599434545),
 ('The R2 score for this alogorithm is :', 0.6849389251211297)}

Hyperparameter Tuning for Random Forest Algorithm

In [75]: `from sklearn.model_selection import RandomizedSearchCV`

```
param_dist = {
    'n_estimators': [50, 100, 200, 300],
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth': [None, 5, 10, 15],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

random_rf = RandomizedSearchCV(estimator = rfr, param_distributions = param_dist, n_iter = 10, cv = 5,
                                scoring = 'neg_mean_squared_error', n_jobs = -1, random_state = 42)
random_rf.fit(x_train, y_train)

# Best Hyperparameter
```

```
best_params = random_rf.best_params_
print(f"Best Hyperparameters: {best_params}")
```

Best Model

```
best_rf = random_rf.best_estimator_
```

Best Hyperparameters: {'n_estimators': 200, 'min_samples_split': 5, 'min_samples_leaf': 2, 'max_features': 'log2', 'max_depth': 15}

In [76]: `evaluation_score(random_rf)`

Out[76]: {('The Adjusted R2 score for this alogorithm is :', 0.7411267383449206),
 ('The Mean Absolute Error for this alogorithm is :', 1292.6247330907584),
 ('The Mean Squared Error for this alogorithm is :', 5314177.19609228),
 ('The R2 score for this alogorithm is :', 0.7421932574004284)}

In [77]: `!pip install xgboost`

```
Requirement already satisfied: xgboost in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (3.0.5)
Requirement already satisfied: numpy in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from xgboost) (2.3.3)
Requirement already satisfied: scipy in c:\users\prerana bochare\appdata\local\programs\python\python313\lib\site-packages (from xgboost) (1.16.2)
[notice] A new release of pip is available: 25.2 -> 25.3
[notice] To update, run: python.exe -m pip install --upgrade pip
```

6. XGBOOST ALGORITHM

In [78]: `import xgboost as xgb`

```
-----  
ModuleNotFoundError                               Traceback (most recent call last)  
Cell In[78], line 1  
----> 1 import xgboots as xgb  
  
ModuleNotFoundError: No module named 'xgboots'
```

In [79]: `import xgboost as xgb`

```
xg_reg = xgb.XGBRegressor(objective = 'reg:squarederror', colsample_bytree = 0.3, learning_rate = 0.1, max_depth = 5,  
                           alpha = 10, n_estimators = 100, random_state = 42)  
  
xg_reg.fit(x_train, y_train)
```

Out[79]:

In [80]: `evaluation_score(xg_reg)`

Out[80]: {('The Adjusted R2 score for this alogorithm is :', 0.7555540795163043),
 ('The Mean Absolute Error for this alogorithm is :', 1406.6248779296875),
 ('The Mean Squared Error for this alogorithm is :', 5018011.5),
 ('The R2 score for this alogorithm is :', 0.7565611600875854)}

Hyperparameter Tuning for XGBoost Algorithm

In [81]:

```
param_dist = {  
    'learning_rate': [0.01, 0.1, 0.2, 0.3],  
    'max_depth': [3, 4, 5, 6],  
    'n_estimators': [50, 100, 200, 300],  
    'subsample': [0.8, 0.9, 1.0],  
    'colsample_bytree': [0.8, 0.9, 1.0],  
    'alpha': [0, 1, 5, 10]  
}  
  
# RandomizedSearchCV for hyperparameter tuning  
  
random_search_xgb = RandomizedSearchCV(estimator = xg_reg, param_distributions = param_dist, n_iter = 10, cv = 5,  
                                         scoring = 'neg_mean_squared_error', n_jobs = -1, random_state = 42)  
random_search_xgb.fit(x_train, y_train)  
  
# Best Hyperparameters  
  
best_params = random_search_xgb.best_params_  
print(f"Best Hyperparameters: {best_params}")  
  
# Best Model  
  
best_xg_reg = random_search_xgb.best_estimator_
```

Best Hyperparameters: {'subsample': 0.8, 'n_estimators': 200, 'max_depth': 3, 'learning_rate': 0.2, 'colsample_bytree': 1.0, 'alpha': 5}

In [82]: `evaluation_score(best_xg_reg)`

Out[82]: {('The Adjusted R2 score for this alogorithm is :', 0.7586664628067963),
 ('The Mean Absolute Error for this alogorithm is :', 1387.6064453125),
 ('The Mean Squared Error for this alogorithm is :', 4954119.5),
 ('The R2 score for this alogorithm is :', 0.7596607208251953)}

7. GRADIENT BOOSTING ALGORITHM

In [83]:

```
from sklearn.ensemble import GradientBoostingRegressor  
  
gb_reg = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)
```

In [84]: `evaluation_score(gb_reg)`

Out[84]: {('The Adjusted R2 score for this alogorithm is :', 0.7262176494071335),
 ('The Mean Absolute Error for this alogorithm is :', 1524.8233966156836),
 ('The Mean Squared Error for this alogorithm is :', 5620232.5219345605),
 ('The R2 score for this alogorithm is :', 0.7273455916754936)}

Hyperparameter Tuning for Gradient Boosting Algorithm

In [85]:

```
from sklearn.model_selection import RandomizedSearchCV  
  
param_dist = {  
    'learning_rate': [0.01, 0.1, 0.2, 0.3],  
    'n_estimators': [50, 100, 200, 300],  
    'max_depth': [3, 4, 5, 6],  
    'min_samples_split': [2, 5, 10],
```

```
'min_samples_leaf': [1, 2, 4],
'subsample': [0.8, 0.9, 1.0],
'max_features': ['auto', 'sqrt', 'log2']
}

random_gbr = RandomizedSearchCV(estimator=gb_reg, param_distributions=param_dist, n_iter=10, cv=5, scoring='neg_mean_squared_error')
random_gbr.fit(x_train, y_train)

best_params = random_gbr.best_params_
print(f"Best Hyperparameters: {best_params}")
best_gb_reg = random_gbr.best_estimator_

Best Hyperparameters: {'subsample': 1.0, 'n_estimators': 200, 'min_samples_split': 10, 'min_samples_leaf': 2, 'max_features': 'sqrt', 'max_depth': 4, 'learning_rate': 0.2}

In [86]: evaluation_score(random_gbr)

Out[86]: {('The Adjusted R2 score for this algorithm is :', 0.7618846187707788),
('The Mean Absolute Error for this algorithm is :', 1330.1576240587872),
('The Mean Squared Error for this algorithm is :', 4888057.2712570755),
('The R2 score for this algorithm is :', 0.7628656184687269)}
```

SUMMARY OF THE PROJECT

Model Comparison Report

The business case of the Flight Fare Data set was to build a machine learning model to help predict the fare of various airlines in the future and let people decide and plan their travel. The data set has 10 features like Airline, Destination, Source, Departure time, Arrival time etc. The data set has a target column (output) as Price of the tickets purchased by the passengers who travelled earlier. Based on these previous records, we had to build a model which could predict the future prices of the tickets and help travelers choose the best priced Airline for their future trips which could be cheaper for them.

To make predictions based on the available data set, different Machine Learning Algorithms were built as mentioned below:

1. Linear Regression
2. Support Vector Machine
3. K-Nearest Neighbour
4. Decision Tree
5. Random Forest
6. XGBoost
7. Gradient Boosting

Out of all models above, **XGBoost Algorithm** and **Gradient Boosting Algorithm** performed best with a **R2 score of 0.75 and 0.73** respectively. After **Hyperparameter Tuning**, all these algorithms gave an improved **R2 score of 0.76**.

In other words, it can be said that 76% of the data are fitting correctly into the model.

Report on Challenges Faced

Missing Values: There were very less missing data in the data set in Route and Total_Stops columns. These null values were dropped as such less missing data will not impact the data set

Date Formatting: The "Date_of_Journey" column was not in a standard date format and this was challenging in performing the date based analyses. Hence, we used the below mentioned technique to handle these data:

- Date Parsing: The "Date_of_Journey" column was converted to a standardized date format (e.g., YYYY-MM-DD) to facilitate date-based calculations and comparisons. This was done by splitting the date string, extracting day, month, and year and then reformatting was done.

Duration Data: The "Duration" column contained values in a non-standard format (e.g., "2h 50m"). Extracting meaningful insights from this column required converting it into a numeric format. The technique used here is as follows:

- Duration Conversion: The "Duration" column was transformed into a numeric format, representing the total duration of the journey in minutes. This involved splitting the string, extracting hours and minutes and converting them into minutes for easy analysis.

Categorical Data: There were several columns like "Airline," "Source," "Destination," and "Total_Stops," that were categorical in nature. These columns were handled using Label Encoder and converted to Numerical Data type.

Conclusion

Using Flight Fare Prediction Dataset, we analyzed various data provided using different exploratory data analysis methods with the help of seaborn and matplotlib libraries. This data set is a supervised learning dataset so we used various regression models to do some predictions about the price of the flight based on some features like type of airline, what is the arrival time, what is the departure time, what is the duration of the flight, source, destination etc.

It can be concluded that among various models applied to this dataset, **XGBoost Algorithm** and **Gradient Boosting Algorithm** performed best. The R2 score for all three algorithms were 0.76 which means that almost 76% of the data fit perfectly into the model and hence do the correct prediction.

In []: